



# FULLSTACK WEB DEV

## mongoose

### BACKEND DEVELOPMENT



MASYNTECH



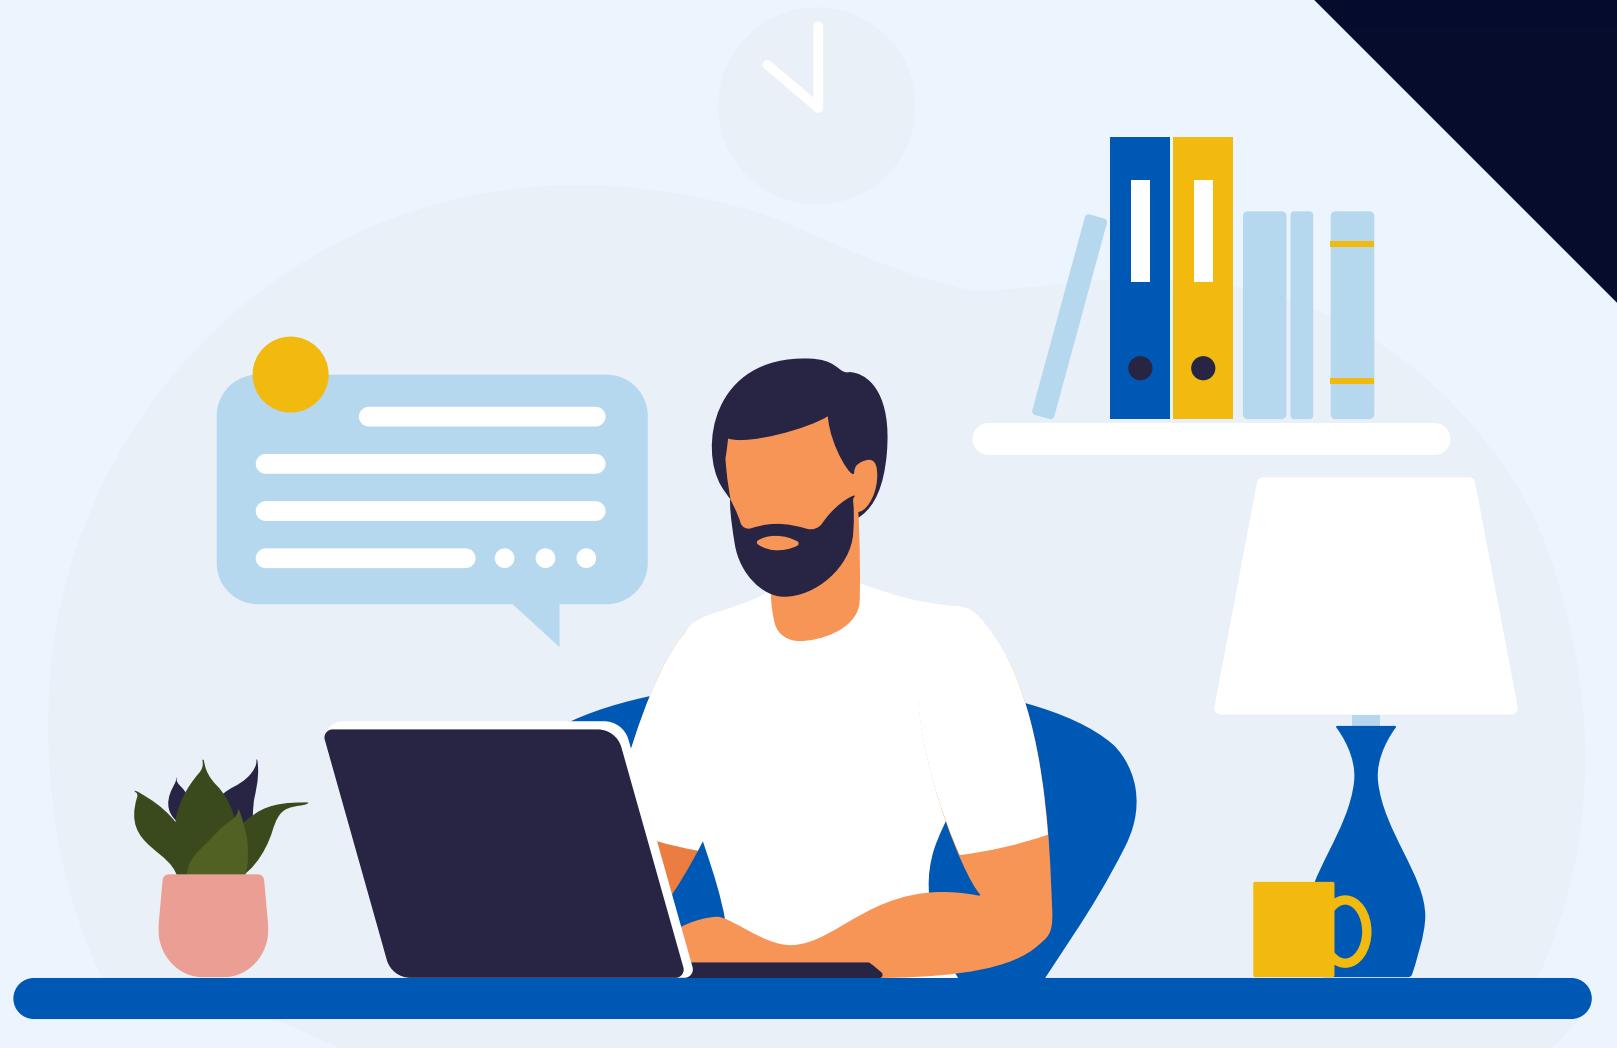
MASYNTECH



[www.masynctech.com](http://www.masynctech.com)



# WHAT'S MONGOOSE



mongoose

CORE CONCEPTS



# MONGOOSE OVERVIEW



## High-Level Explanation

- Mongoose: ODM library for MongoDB in Node.js
- Offers schema-based models, validation, queries, hooks, etc.
- Simplifies MongoDB interaction in Node.js apps.



## Deep Dive

- Mongoose abstracts MongoDB for ease
- Handles validation, casting, queries, and hooks
- Defines document structure with Schemas
- Features: population, virtual properties, middleware, plugins
- Maintains MongoDB's flexibility with added structure



## Analogue

- Analogy: Mongoose as a toy box organizer
- Keeps MongoDB data neat and organized
- Ensures data matches your desired structure



## When to use?

- Mongoose benefits for:
  - Complex projects with structured DB needs
  - Ensuring data adheres to rules/schemas
  - Quick prototyping for speed and convenience



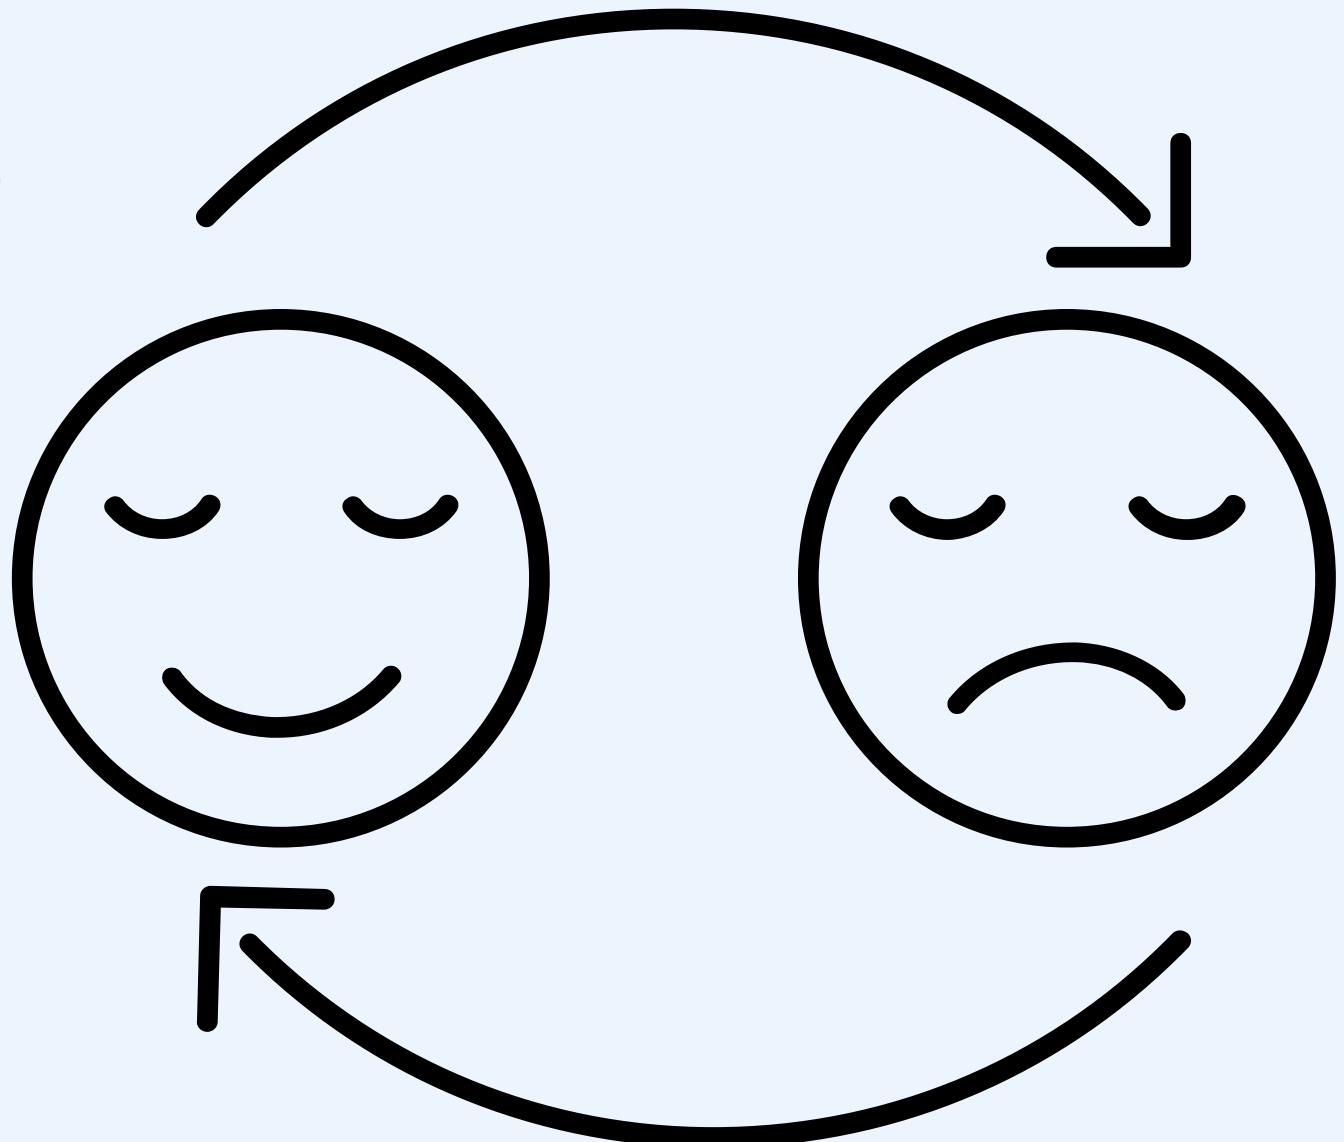
## Best Practices

- Steps for Mongoose usage:
  - Define schema for collections
  - Employ validators for data integrity
  - Utilize middleware for pre/post-save operations



# MONGODB VS MONGOOSE

mongoose



CORE CONCEPTS



# MONGOOSE VS MONGODB



## High-Level Explanation

- Mongoose vs. MongoDB native driver:
  - Mongoose: Higher-level, schema-based, validation
  - MongoDB native driver: Lower-level, direct access



## Deep Dive

- **Mongoose:**
  - ODM layer with schemas, rich features.
  - Strongly-typed schemas, hooks, virtuals.
  - Eases development but adds some abstraction.



## Analogue

- Mongoose: Organized kitchen with helpers for correct data storage.
- MongoDB Native Driver: Basic kitchen without organizers; user responsibility for data handling.



## When to use?

- **Mongoose:**
  - Leverage schemas and built-in validation.

- **MongoDB Native Driver:**
  - Manage data integrity manually, exercise caution in DB operations.

- **Mongoose:**
  - Complex apps, data validation, ease of use.
  - Speedy development.

- **MongoDB Native Driver:**
  - Simple CRUD, intricate MongoDB operations.



# MONGOOSE OFFICIAL DOCS

`mongoose`



CORE CONCEPTS



# MONGOOSE

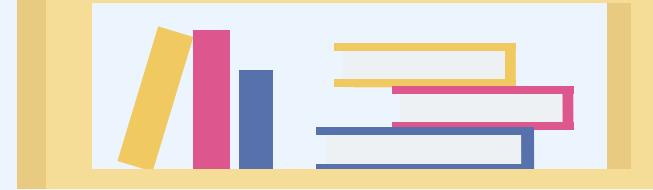
## INSTALLATION

mongoose



CORE CONCEPTS

# CONNECT TO MONGOOSE



mongoose

CORE CONCEPTS



# SCHEMA OVERVIEW



mongoose

CORE CONCEPTS



# SCHEMA OVERVIEW



## High-Level Explanation

- MongoDB schema: Blueprint for collection documents
- Mongoose schema: Defines fields, types, validators, structure



## Deep Dive

- **Mongoose schemas:**
  - Define field types, default values, validation
  - Create virtual properties for app logic
  - Maintain data integrity in schema-less MongoDB
  - Manage complex relationships, validation, hooks, methods



## Analogue

- Analogy: MongoDB is like open LEGO pieces.
- Mongoose is the rulebook for building a stable LEGO castle.
- Ensures walls, towers, drawbridge follow the right design.



## When to use?

### Use Mongoose schemas when:

- Structured data with specific types needed.
- Data validation before database save crucial.
- Pre- or post-save hooks required.
- Handling complex relationships between documents.



## Best Practices

- Define document shape with Mongoose schema.
- Apply suitable field validation.
- Minimize use of `any` data type.
- Implement pre- and post-save hooks as needed.



# SCHEMA TYPES



mongoose

CORE CONCEPTS



# SCHEMA TYPES



## High-Level Explanation

- Mongoose schema types:
  - `String`, `Number`, `Date`, `Boolean`, `Array`, `ObjectID`
- Used to define structure of MongoDB documents.



## Analogue

- Analogy: Schema types as LEGO blocks
- `String`: Long, flat piece for names, descriptions
- `Number`: Small, square block for age, price
- `Array`: Tray for multiple blocks of same/different shapes



## Best Practices

- Opt for the most restrictive type for data integrity.
- Minimize use of `Mixed` unless compelling need.
- Think about performance implications.
- Use `ObjectID` for document references in relational data.



## Deep Dive

- Mongoose schema types:
  - `String`: Textual data, symbols, numbers.
  - `Number`: Integers and floats.
  - `Date`: Date and optional time.
  - `Boolean`: True or false.
  - `Array`: Multiple values, often of the same type.
  - `ObjectID`: Unique identifier, typically auto-generated.
  - `Mixed`: Flexible, holds any data type (use sparingly).
- Embedded Documents: Nested, complex structures within schemas.



## When to use?

- Use specific schema types when you:
- Have precise data expectations.
  - Need validation or transformation.
  - Want data integrity assurance.
  - Seek query performance optimization.



# DEMO TIME



# MODELS

# OVERVIEW

mongoose



CORE CONCEPTS



# MODELS OVERVIEW



## High-Level Explanation

- Mongoose model: Interface for MongoDB interactions.
- Created by combining schema with CRUD methods.
- Defines fields, types, and how to work with data.



## Analogue

- Analogy: Model as a blueprint for houses
- Blueprint (Schema): Defines walls, doors, windows (fields)
- Build houses (documents) based on blueprint
- Find houses with specific features (querying)



## Best Practices

- **Single Responsibility:** One model, one concept.
- **Naming:** Singular, capitalized model names.
- **Method Attachment:** Attach custom logic via instance or static methods.



## Deep Dive

- Models construct MongoDB documents.
- Compiled from a schema.
- Map to MongoDB collections.
- Define document structure.
- Perform CRUD and aggregations.
- Components: Schema, Collection, Document, Query Interface.

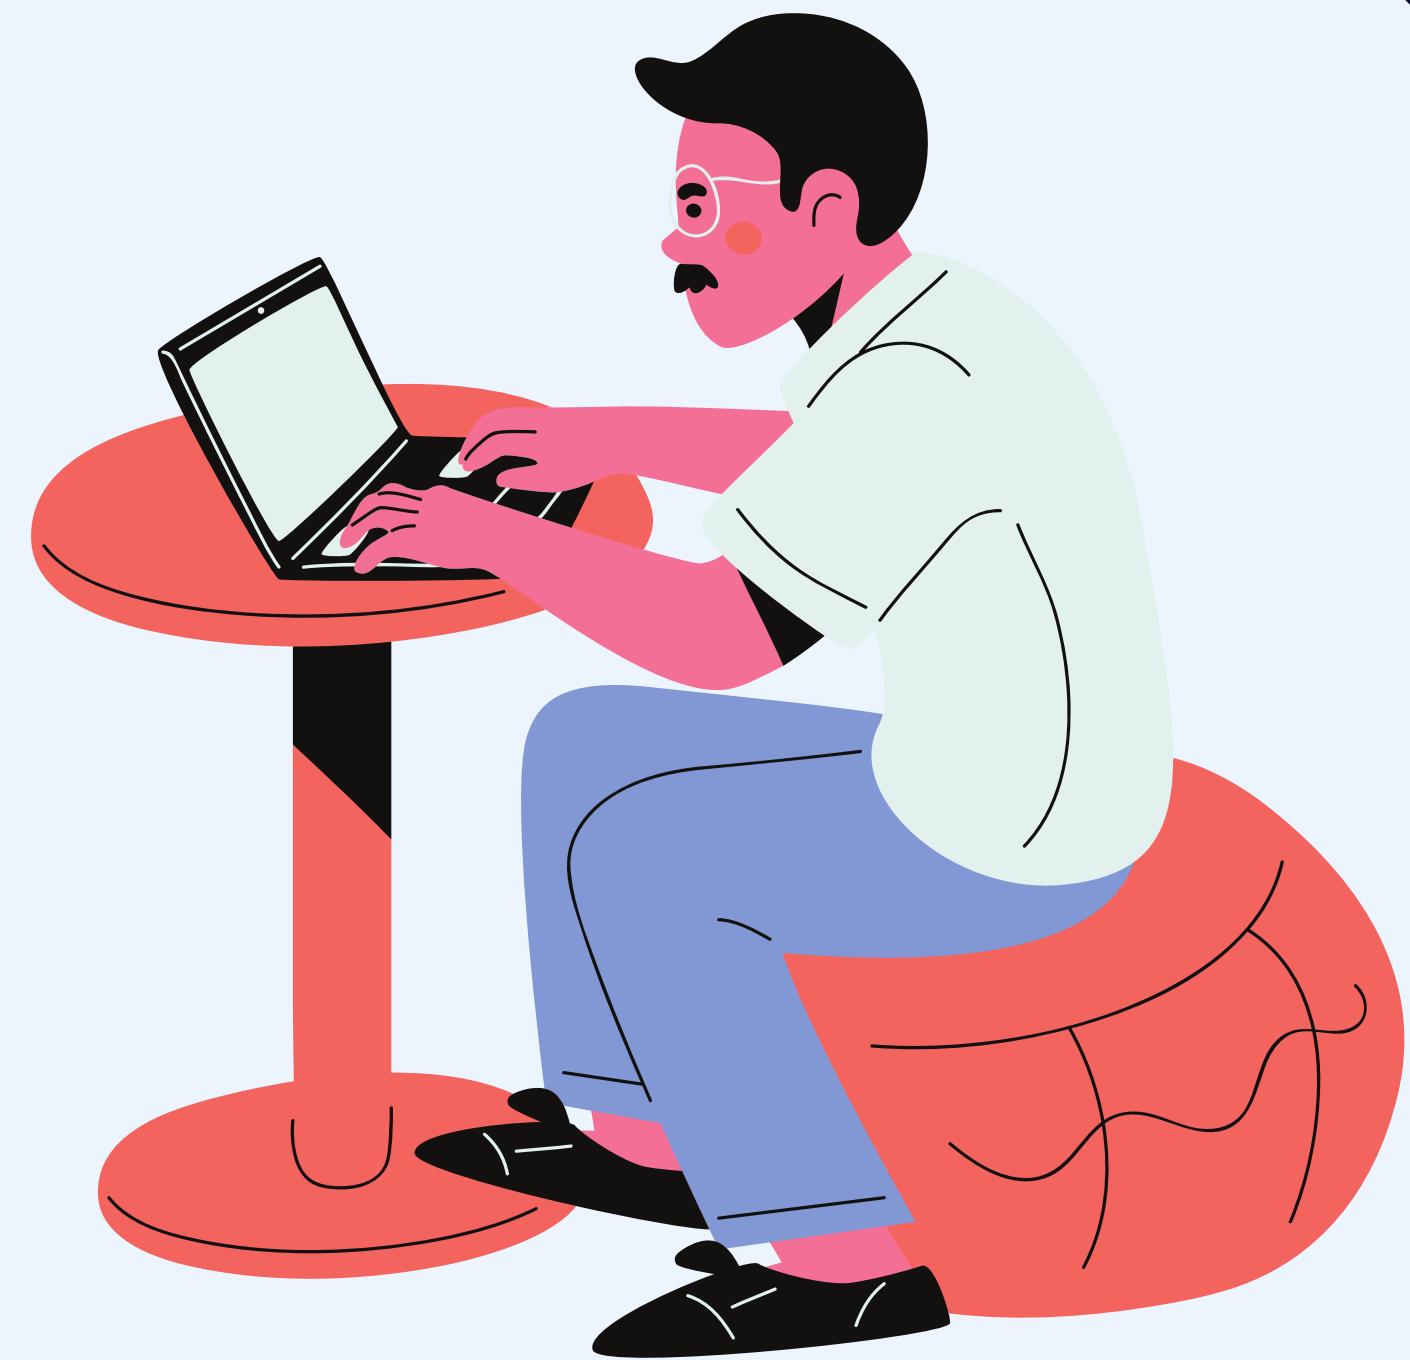


## When to use?

- Use models when you:
- Perform read/write operations in MongoDB.
  - Need to enforce structure on collection documents.
  - Want to encapsulate data-related business logic (preferably simple logic).



# DEMO TIME





# MONGOOSE CRUD OPERATIONS



mongoose

# CREATE OPERATIONS (`.save()`)



mongoose

CRUD OPERATIONS



# CREATE OPERATIONS USING `SAVE()



## High-Level Explanation

- `save()` in Mongoose: Create operation, saves new document to DB.



## Analogue

- Analogy: `save()` in Mongoose like 'Save Game' in video games.
- Stores the new 'game state' (document) in MongoDB.



## Best Practices

- Always handle promises (`.then/.catch` or `async/await`).
- Validate input before `save()` to prevent incorrect data.
- Prepare for possible failures like duplicate keys.



## Deep Dive

- Instantiate a Mongoose object using a model for an instance.
- Invoke `save()` on the instance.
- Mongoose handles validation, transformation, MongoDB `insert`.
- Promise resolves on success, rejects on failure.
- Catch errors with `catch()`.



## When to use?

### Use `save()` when:

- Inserting a new document into the collection.
- Updating an existing document after making changes.



# DEMO TIME





# CREATE OPERATIONS

`(.create())`



mongoose

CRUD OPERATIONS

# CREATE OPERATIONS USING THE `CREATE()``



## High-Level Explanation

- `.create()` in Mongoose: Shortcut for saving one or more documents to MongoDB.
- Combines model instantiation and `.save()` in one step.



## Analogue

- Analogy: Traditional vs. `.create()` method in video game.
- Traditional: Gather materials, go to site, build.
- `.create()`: Cheat code, instant house construction.



## Best Practices

- Handle errors with the returned promise (`.then` or `await`).
- Validate data before using `.create()`; schema validations run.
- Caution with multiple document insertions; one failure blocks all.



## Deep Dive

- Traditionally: Instantiate, modify, and `.save()` for a new document.
- `.create()` simplifies, creating and saving in one step.
- Makes code more concise and maintainable.



## When to use?

- Use `.create()` when:
  - Swift document insertion without prior manipulation.
  - Inserting multiple documents with confidence.
  - Aiming for concise and clean code.





# DEMO TIME





# CREATE OPERATIONS

`(.insertMany())`



mongoose

CRUD OPERATIONS

# BATCH OPERATIONS USING `\$.INSERTMANY()



## High-Level Explanation

- `\$.insertMany()` in MongoDB: Insert multiple docs in one operation.
- Useful for batch processing, efficient data insertion.



## Analogue

- Analogy: `\$.insertMany()` like dumping Lego bricks into a big box.
- Saves time and effort by inserting multiple docs at once.



## Best Practices

- Validate documents against schema.
- Watch document size limit (16MB).



## Deep Dive

- `\$.insertMany()` : Insert multiple items in collection at once.
- Ideal for seeding databases or efficient bulk inserts.



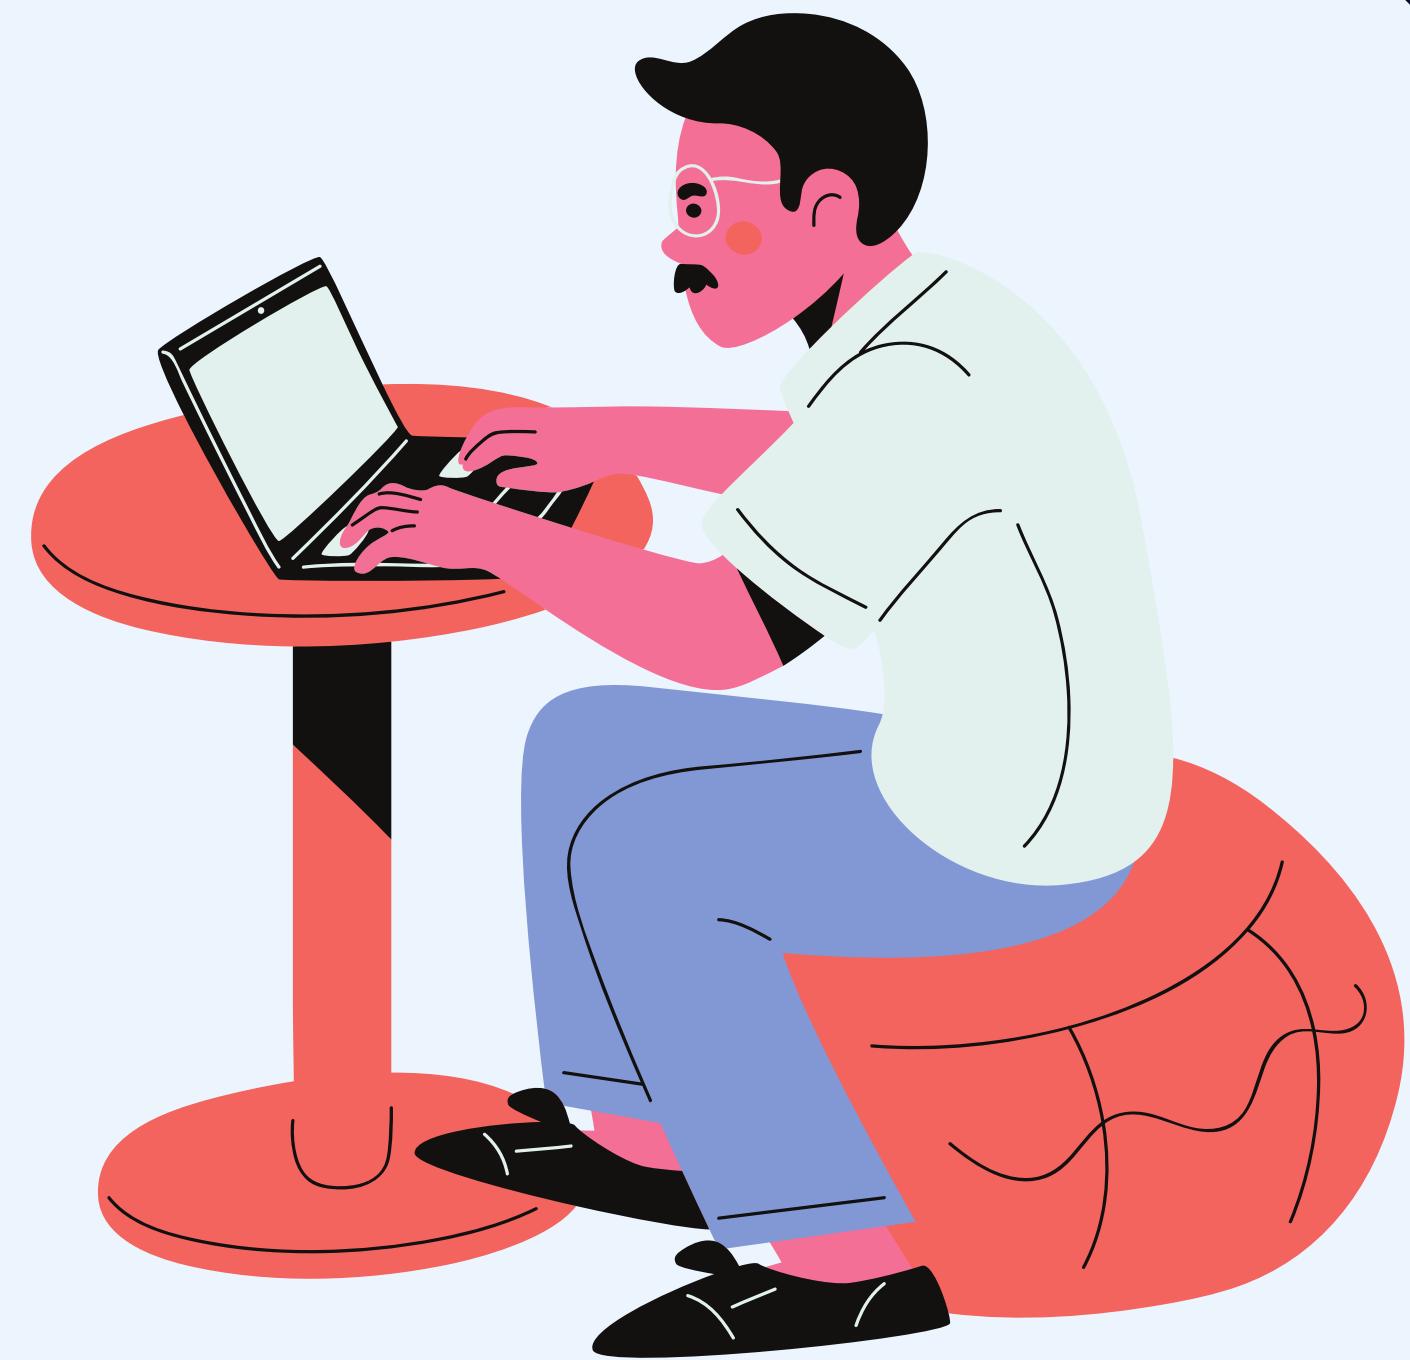
## When to use?

- Use `\$.insertMany()` when:
- Seeding a database with initial data.
  - Inserting an array of objects into a collection.





# DEMO TIME





# CREATE OPERATIONS

`(.insertMany())`



mongoose

CRUD OPERATIONS

# BATCH OPERATIONS USING `\$.INSERTMANY()



## High-Level Explanation

- `\$.insertMany()` in MongoDB: Insert multiple docs in one operation.
- Useful for batch processing, efficient data insertion.



## Deep Dive

- `\$.insertMany()` : Insert multiple items in collection at once.
- Ideal for seeding databases or efficient bulk inserts.



## Analogue

- Analogy: `\$.insertMany()` like dumping Lego bricks into a big box.
- Saves time and effort by inserting multiple docs at once.



## When to use?

- Use `\$.insertMany()` when:
- Seeding a database with initial data.
  - Inserting an array of objects into a collection.

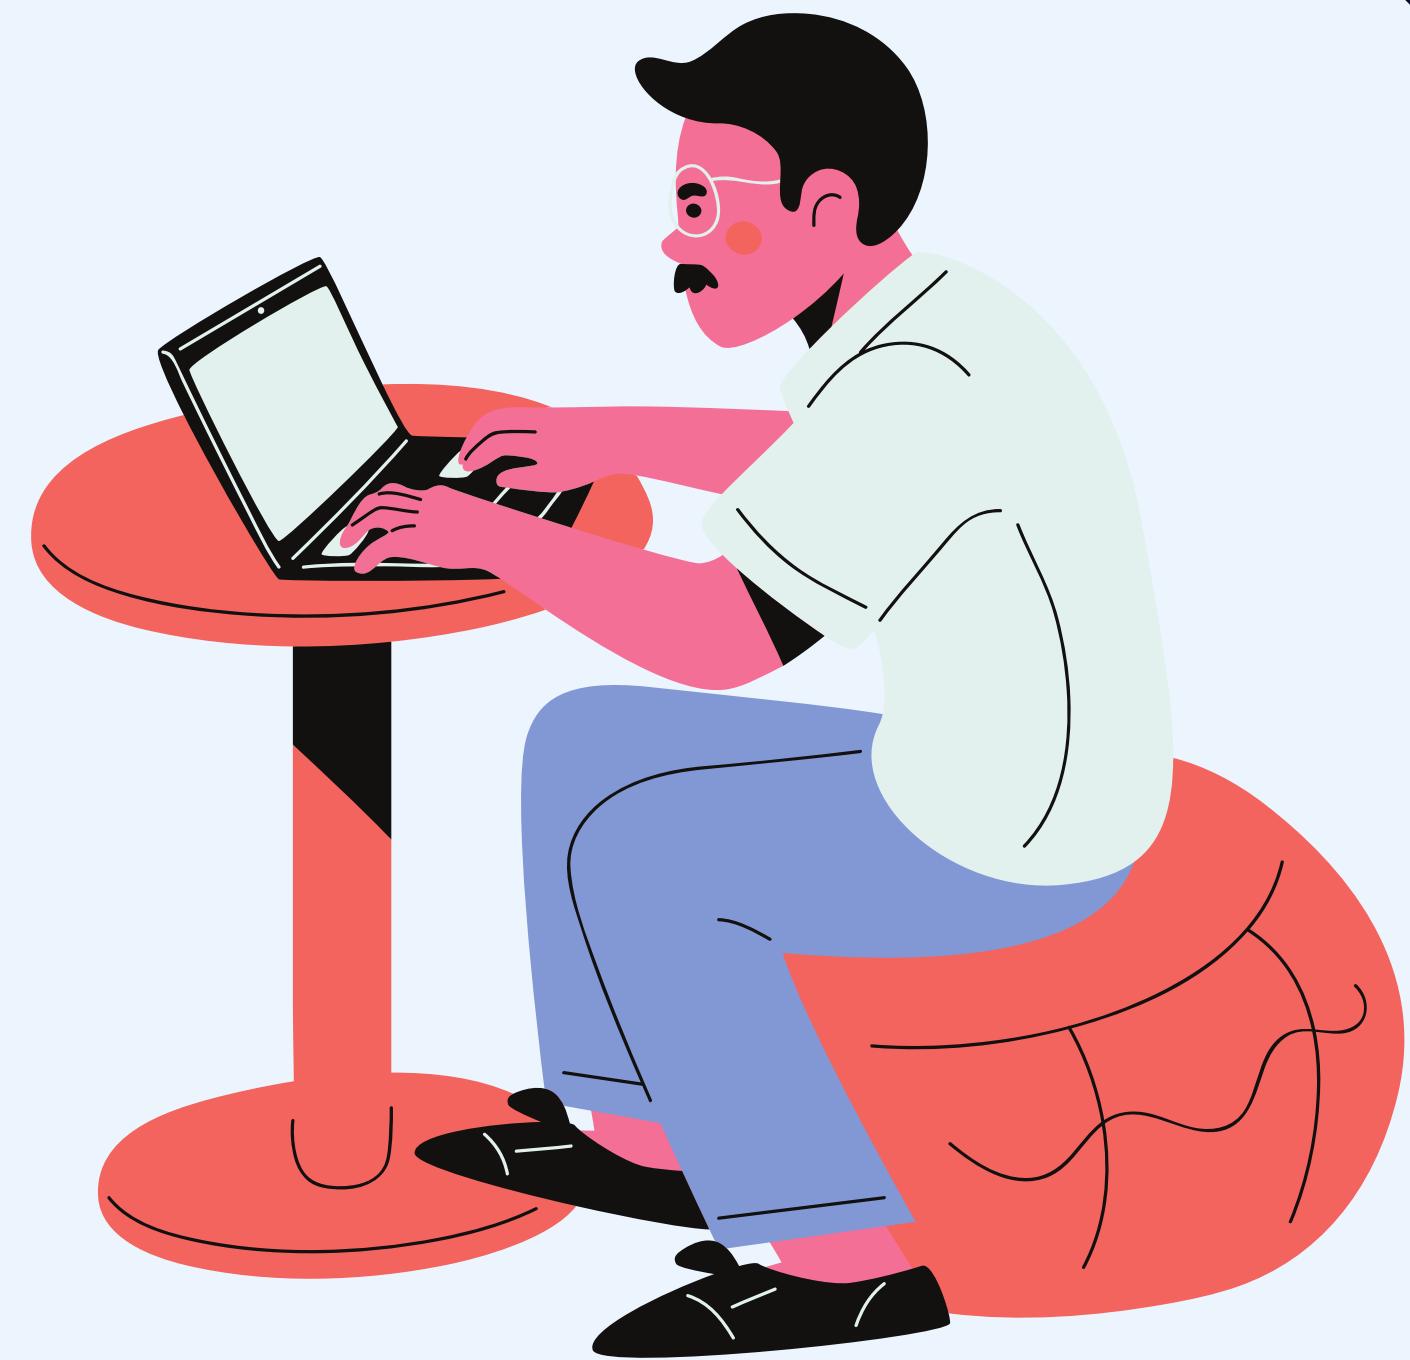


## Best Practices

- Validate documents against schema.
- Watch document size limit (16MB).



# DEMO TIME





# READ OPERATIONS



mongoose

CRUD OPERATIONS

# READ OPERATIONS .FIND() AND .FINDONE()



## High-Level Explanation

- `.find()`: Fetch multiple documents based on query.
- `.findOne()`: Retrieve single document based on criteria.



## Analogue

- Analogy: `.find()` like scrolling through matches.
- `.findOne()` like picking the top match based on preferences.



## Deep Dive

- `.find()`: Returns array of matching docs or empty array.
- `.findOne()`: Returns first matching doc or `null`.
- Both accept projections, sorting, limiting, and skipping.



## When to use?

- `.find()`: For multiple docs, like all products in a category.
- `.findOne()`: For a single doc, like user details by username.



## Best Practices

- Projections: Limit fields returned, reduce data transfer.
- Options: `sort`, `limit`, `skip` for query fine-tuning.
- Handle errors using try-catch or `.catch()`.



# DEMO TIME





# READ QUERIES OPERATIONS



mongoose

QUERY OPERATIONS

# READ OPERATIONS `WHERE()`, `SORT()`, `LIMIT()`



## High-Level Explanation

- **.where()**, **.sort()**, and **.limit()** refine query results.
- Typically chained to `.find()` or `.findOne()` queries.



## Analogue

- **.where()**: "Red" filter
- **.sort()**: "Sort by Price: Low to High"
- **.limit()**: "Show only the first 5 options"



## Best Practices

- **.where()**: Validate input to prevent injection attacks.
- **.sort()**: Let the database handle sorting for efficiency.
- **.limit()**: Be cautious with the limit to avoid missing data or overwhelming the system.



## Deep Dive

- **.where()**: Filter docs with conditions, logical operators.
- **.sort()**: Sort docs based on fields.
- **.limit()**: Limit the number of returned docs.



## When to use?

- **.where()**: Complex filtering with multiple conditions.
- **.sort()**: Organizing data for readability or analysis.
- **.limit()**: Pagination or performance optimization by limiting records.



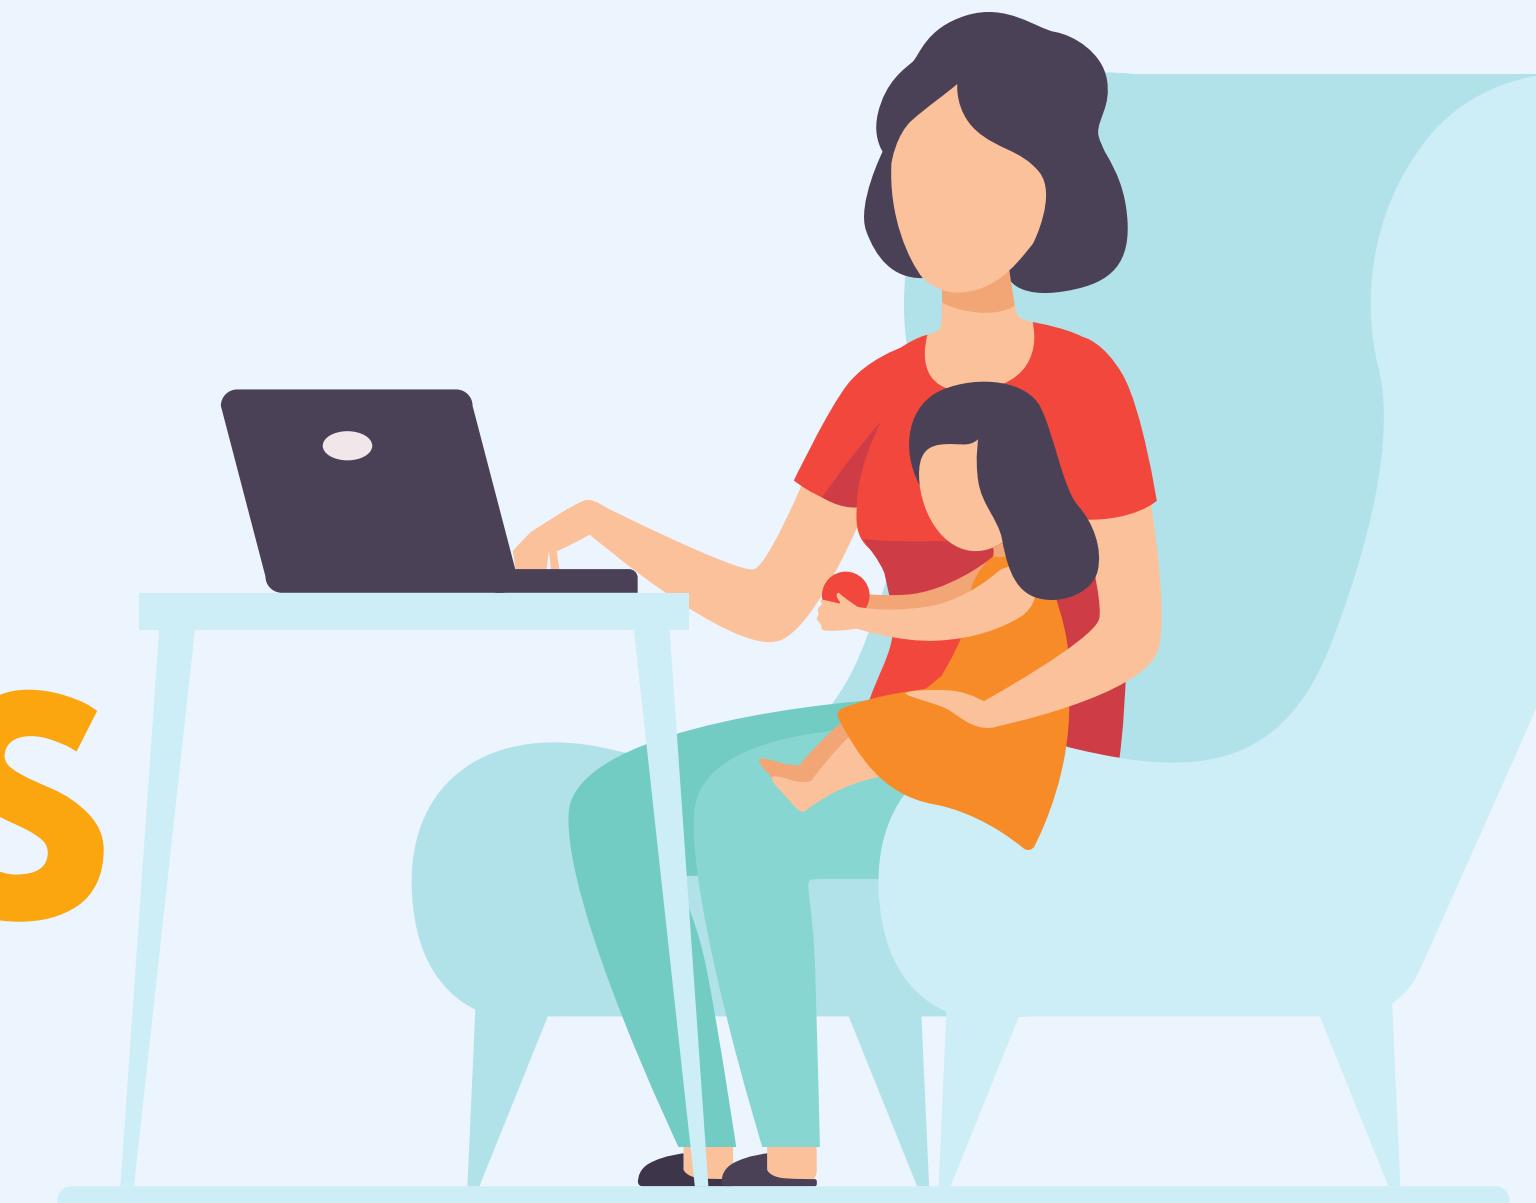
# DEMO TIME



# READ QUERIES OPERATIONS

## Filtering Results

mongoose



QUERY OPERATIONS



# READ OPERATIONS USING QUERY CONDITIONS



## High-Level Explanation

- MongoDB query conditions filter returned documents.
- Operators like ` `\$eq` , ` \$gt` , ` \$lt` , ` \$in` , ` \$nin` used for criteria.
- Essential for precise search results.



## Analogue

- Analogous to music playlist filters
- ` `\$eq` : Genre equals a certain type
- ` `\$gt` : Released after 2010
- ` `\$nin` : Not in disliked genres list
- Filters tailor playlist to specific taste



## Best Practices

- Optimize by indexing frequently queried fields.
- Ensure data validation to prevent misuse.
- Enhance readability with helper methods (e.g., ` .gte()` , ` .lt()` ).



## Deep Dive

- Query conditions: Key-value pairs
- Key: Field of interest
- Value: Condition (often another key-value pair)
- Examples:
  - ` `\$eq` : Equal
  - ` `\$gt` : Greater than
  - ` `\$lt` : Less than
  - ` `\$in` : Matches any in array
  - ` `\$nin` : Matches none in array



## When to use?

- Query conditions ideal for large datasets
- Valuable in:
  - Data analytics
  - Reporting tools
  - Rapid info retrieval in real-time apps



# DEMO TIME



# QUERY PROJECTION

## Selecting fields

mongoose



QUERY OPERATIONS

# PROJECTIONS: SELECTING SPECIFIC FIELDS



## High-Level Explanation

- MongoDB projections: Select specific document fields
- Include/exclude fields, boost query performance
- Minimize data transfer over network



## Analogue

- Analogous to phone contact list
- Display only necessary info (name, photo)
- Like MongoDB projections, faster access



## Best Practices

- Avoid mixing inclusion/exclusion, except for `\\_id`
- Use projections carefully for efficiency
- Be mindful of array field results



## Deep Dive

- MongoDB projection in `find()` and `findOne()`
- Object specifies inclusion (`1` or `true`) or exclusion (`0` or `false`)
- Exception: `\\_id` field included by default, can be explicitly excluded.



## When to use?

- Reasons to use MongoDB projections:
  - Minimize network data
  - Crucial for data-sensitive mobile apps
  - Tailor fields for specific use cases



# DEMO TIME



# UPDATE OPERATIONS

Using `.**update()**` and `.**updateOne()**...



mongoose

QUERY OPERATIONS



# UPDATE OPERATIONS



## High-Level Explanation

- **updateOne**: Update a single document.
- **updateMany**: Update multiple documents.
- **findByIdAndUpdate**: Find by ID and update a document.
- **findOneAndUpdate**: Find one document and update it.



## Analogue

- Analogy: Twitter feed as MongoDB collection, tweets as documents
- **findOneAndUpdate**: Modify single tweet with conditions
- **updateMany**: Change tweets with specific hashtag
- **findByIdAndUpdate**: Update tweet with known ID



## Best Practices

- Validate updates to avoid unintended changes.
- Use `{ new: true }` option for updated document return.
- Note: Not all update methods trigger Mongoose middleware.



## Deep Dive

- Use appropriate Mongoose update method for operation scale.
- Define update conditions and changes.
- Framework interacts with MongoDB for updates.
- Account for update options, default values, and middleware.



## When to use?

- **updateOne**: Single doc, specific criteria.
- **updateMany**: Bulk updates, multiple matching criteria.
- **findByIdAndUpdate**: Unique ID known.
- **findOneAndUpdate**: Update first matching filter.



# DEMO TIME



# UPDATE OPERATORS

mongoose



QUERY OPERATIONS



# UPDATE OPERATORS



## High-Level Explanation

- Mongoose update operators modify field values in documents.
- Useful for changing values without replacing the entire doc.



## Deep Dive

### Field Update Operators:

- `'\$set` : Set field value in document.
- `'\$unset` : Remove field from document.

### Array Update Operators:

- `'\$addToSet` : Add item if not exists in array.
- `'\$push` : Add item to array.
- `'\$pop` : Remove item from array.
- `'\$pull` : Remove all instances of value from array.
- `'\$pullAll` : Remove all matching values from array.

### Arithmetic Operators:

- `'\$inc` : Increment field's value.
- `'\$mul` : Multiply field's value.
- `'\$min` : Update field to minimum of current and given values.
- `'\$max` : Update field to maximum of current and given values.

### Date Operators:

- `'\$currentDate` : Set field to current date.



# DEMO TIME



# DELETE OPERATIONS

Using ``.remove()``, ``.deleteOne()``, and ``.deleteMany()``



mongoose

QUERY OPERATIONS



# DELETE OPERATIONS



## High-Level Explanation

- **`deleteOne`**: Delete a single document based on criteria.
- **`deleteMany`**: Remove multiple documents matching criteria.
- **`findByIdAndDelete`**: Find by ID and delete a document.
- **`findOneAndDelete`**: Find one document and delete it.



## Analogue

- Analogy: Collection as deck of cards
- **`deleteOne`**: Remove first 'Heart'
- **`deleteMany`**: Remove all 'Hearts'
- **`findByIdAndDelete`**: Remove card by ID
- **`findOneAndDelete`**: Remove and show first 'Heart'



## Deep Dive

- **`deleteOne`**: Removes first matching document by filter.
- **`deleteMany`**: Deletes all matching documents by filter.
- **`findByIdAndDelete`**: Deletes document by unique `\\_id`.
- **`findOneAndDelete`**: Deletes and returns first matching doc by query.



## When to use?

- **`deleteOne`**: Remove first matching doc.
- **`deleteMany`**: Batch deletion with common attribute.
- **`findByIdAndDelete`**: Delete by known `\\_id`.
- **`findOneAndDelete`**: Show data of deleted doc.



## Best Practices

- Test deletions on dev database first.
- Use callbacks/promises for error handling or post-deletion actions.
- Be cautious with `deleteMany` in production.

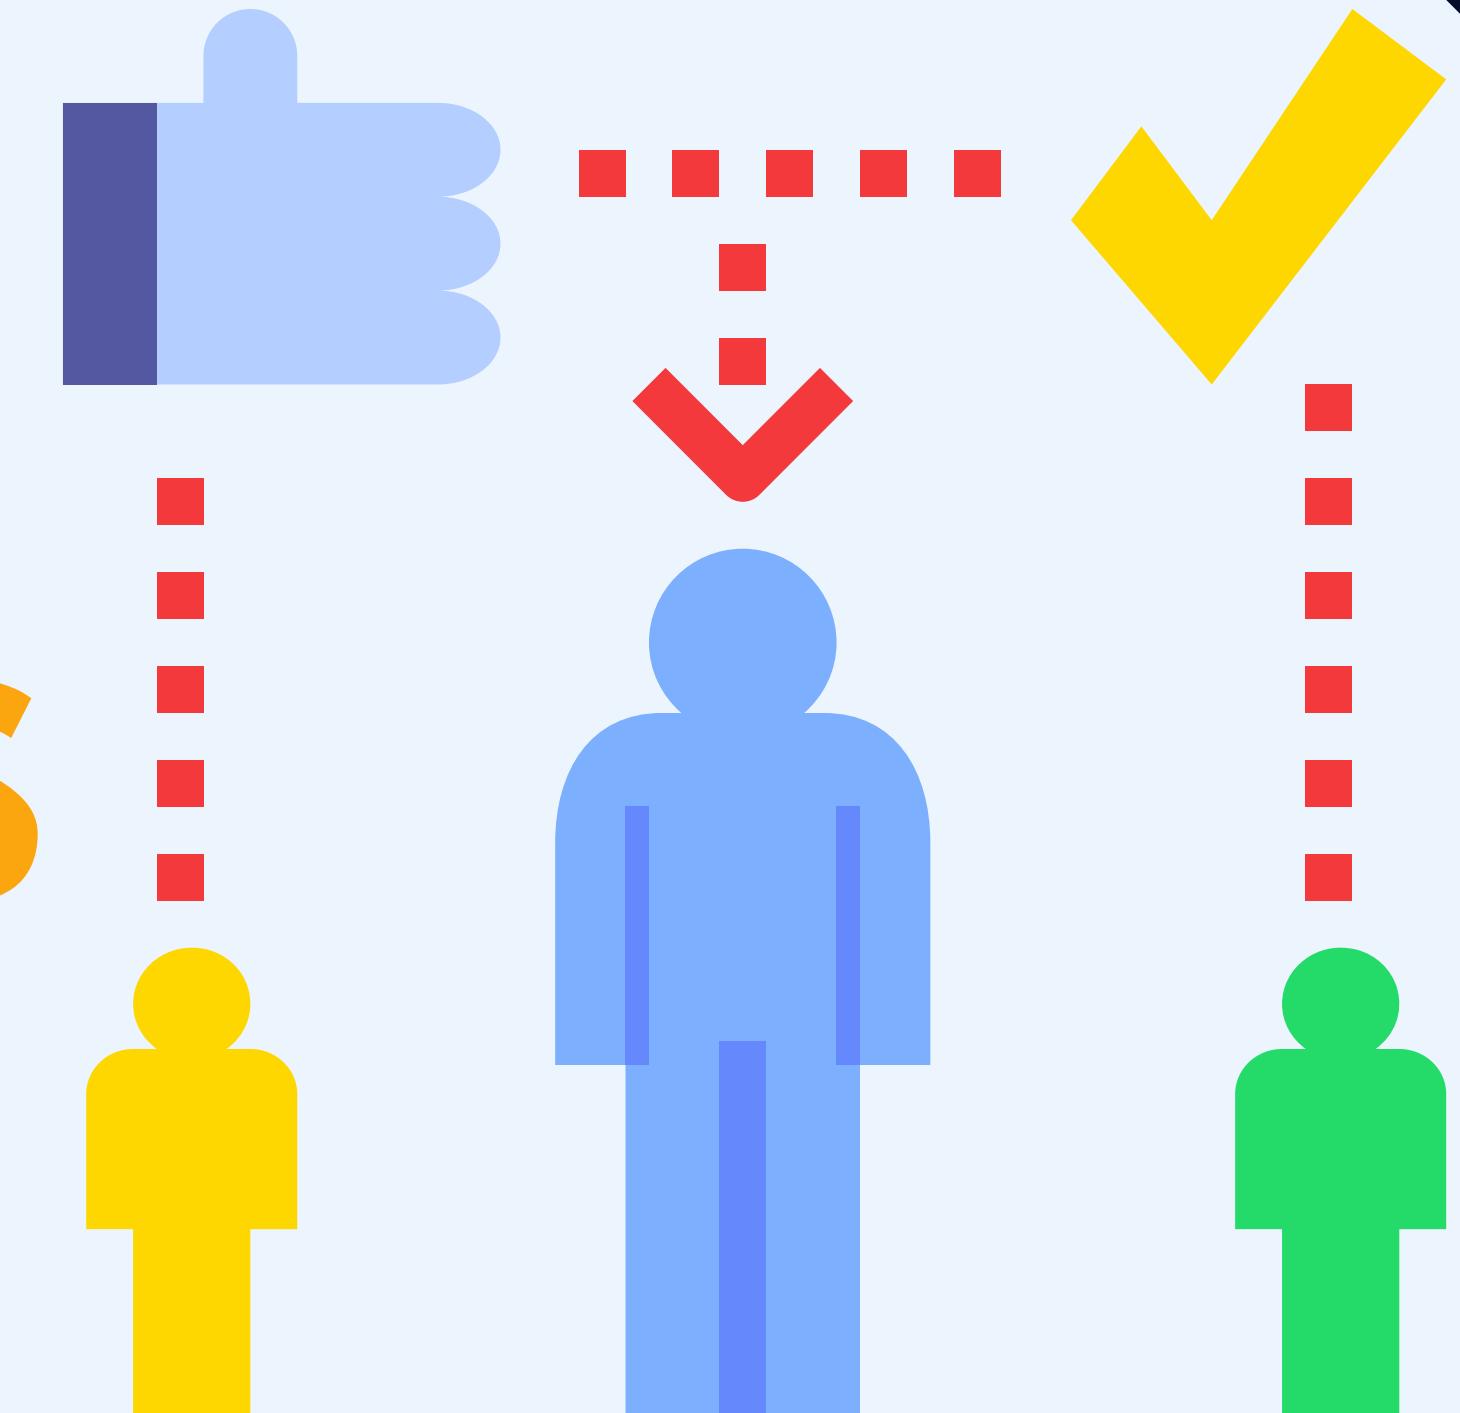


# DEMO TIME





# VALIDATIONS

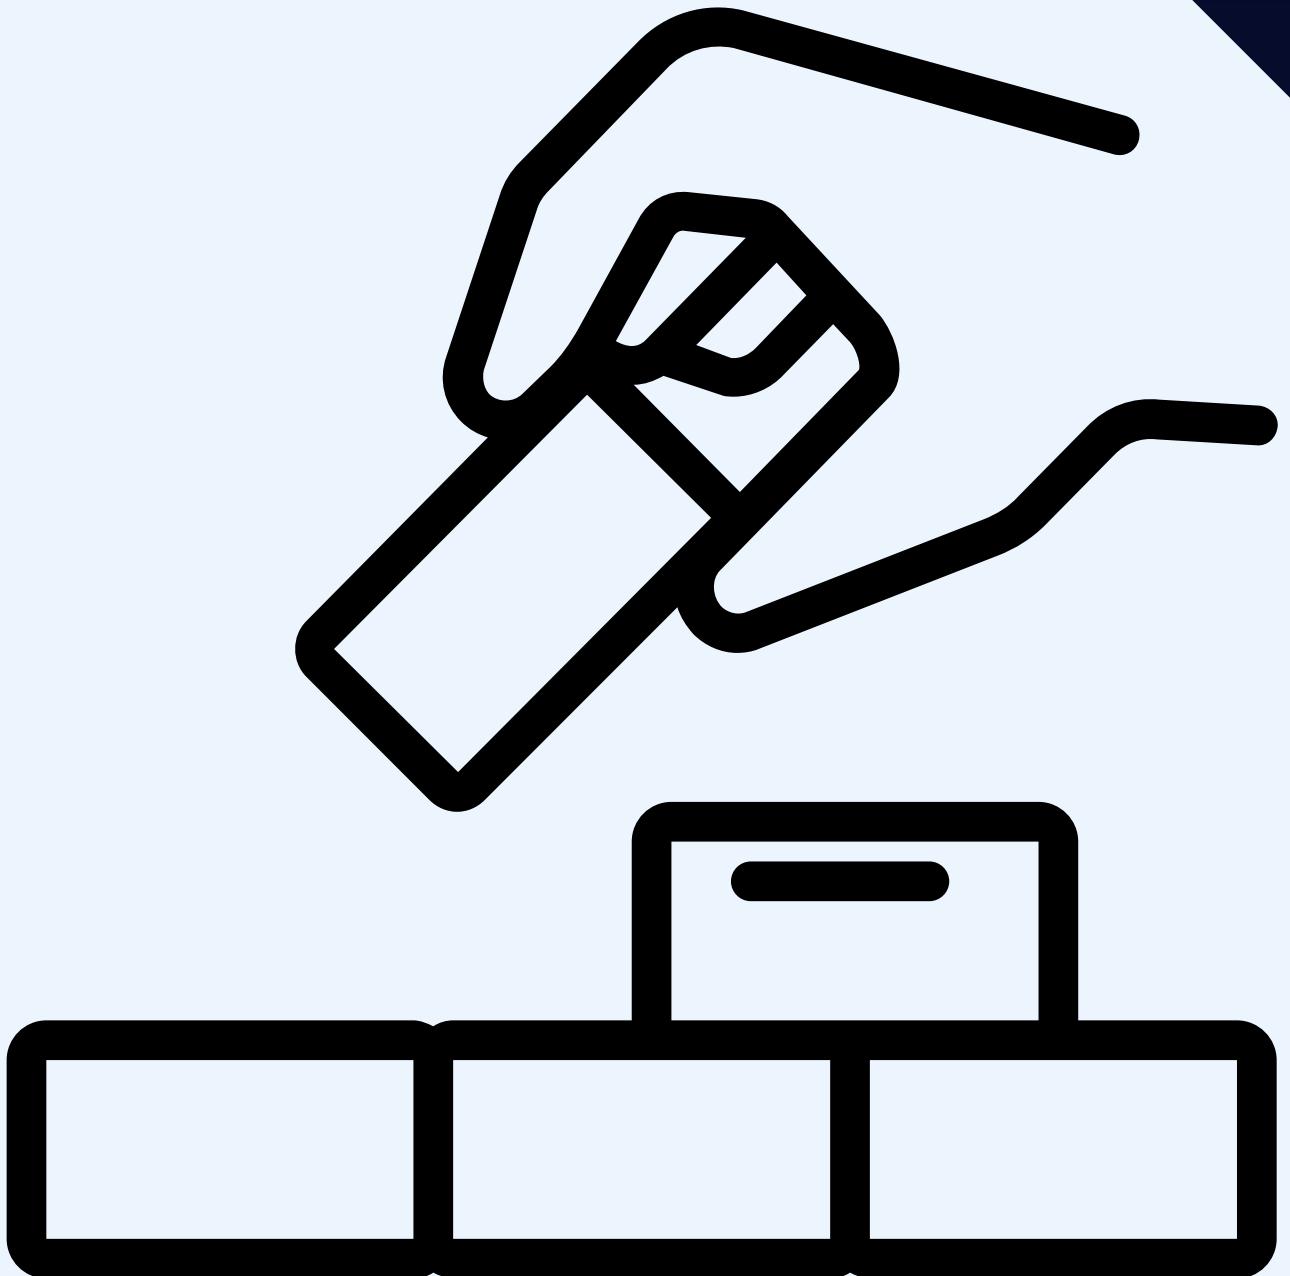


mongoose

MONGOOSE VALIDATION



# BUILT-IN VALIDATORS



mongoose

MONGOOSE VALIDATION



# BUILT-IN MONGOOSE VALIDATORS



## High-Level Explanation

- Mongoose validators: Enforce data validation on MongoDB docs.
- Predefined checks like `required`, `min`, `max`, `enum`, `unique`, etc.
- Streamline data validation process.



## Analogue

- Analogy: Mongoose validators like spell checker.
- Built-in features prevent mistakes.
- E.g., field for email only, Mongoose enforces it.



## Best Practices

- Pair with server-side checks for robust validation.
- Don't rely solely on `unique` for critical uniqueness.
- Provide meaningful error messages for client-side feedback.



## Deep Dive

**Required:** Field must exist with a value.

**Type:** Enforce datatype (String, Number, Date, etc.).

**Minlength/Maxlength:** Set string length limits.

**Validate:** Custom async validation.

**Min/Max:** Enforce numerical range.

**Enum:** Value must match options in an array.

**Unique:** Ensure uniqueness within a collection (schema option).

**Match:** Validate against a regex pattern.



## When to use?

- Use Mongoose built-in validators for:
  - Quick, efficient validation.
  - Standard requirements without custom logic.

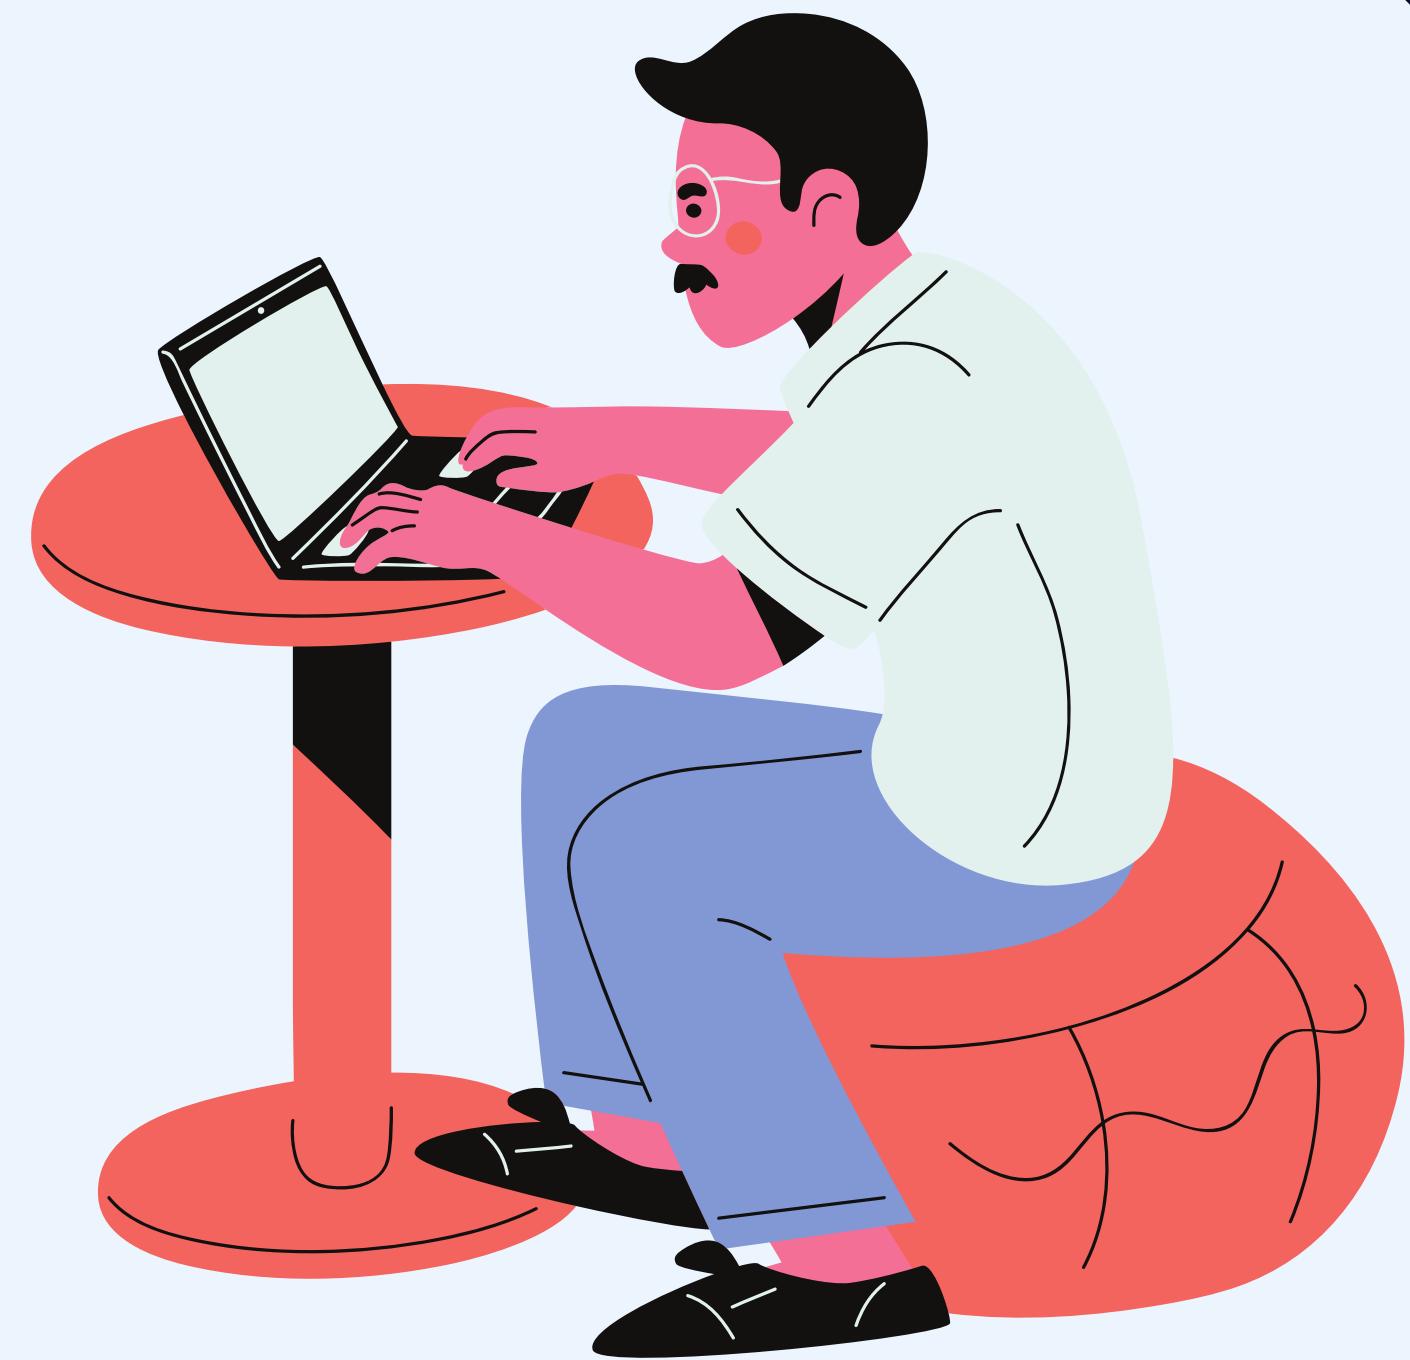


## Summary

- Mongoose built-in validators: Quick, reliable, standard rules.
- Cover field requirements, types, value constraints.
- Maintain data integrity without extensive custom code.



# DEMO TIME





# CUSTOM VALIDATION LOGIC

mongoose



MONGOOSE VALIDATION



# CUSTOM VALIDATION LOGIC



## High-Level Explanation

- Custom validation logic in Mongoose: Define own rules.
- Use `validate` property in schema definition.
- Implement custom function returning `true` or `false` for validation.



## Analogue

- Analogy: Hosting a party, acting as a custom validator
- Players pick specific balls based on custom rules
- Ensuring everyone follows the unique criteria



## Best Practices

- Handle both synchronous and asynchronous errors.
- Provide clear and meaningful error messages.
- Rigorously test your custom validation logic.
- Keep validation functions as lightweight as possible.



## Deep Dive

- Mongoose's `validate` attribute: Custom validation function
- Returns `true` for valid, `false` for invalid
- Supports `Promise` for async validation
- Use for custom criteria beyond built-in validators



## When to use?

- Built-in validators are inadequate.
- Complex logic spans multiple fields.
- Asynchronous operations, like API calls, are required for validation.

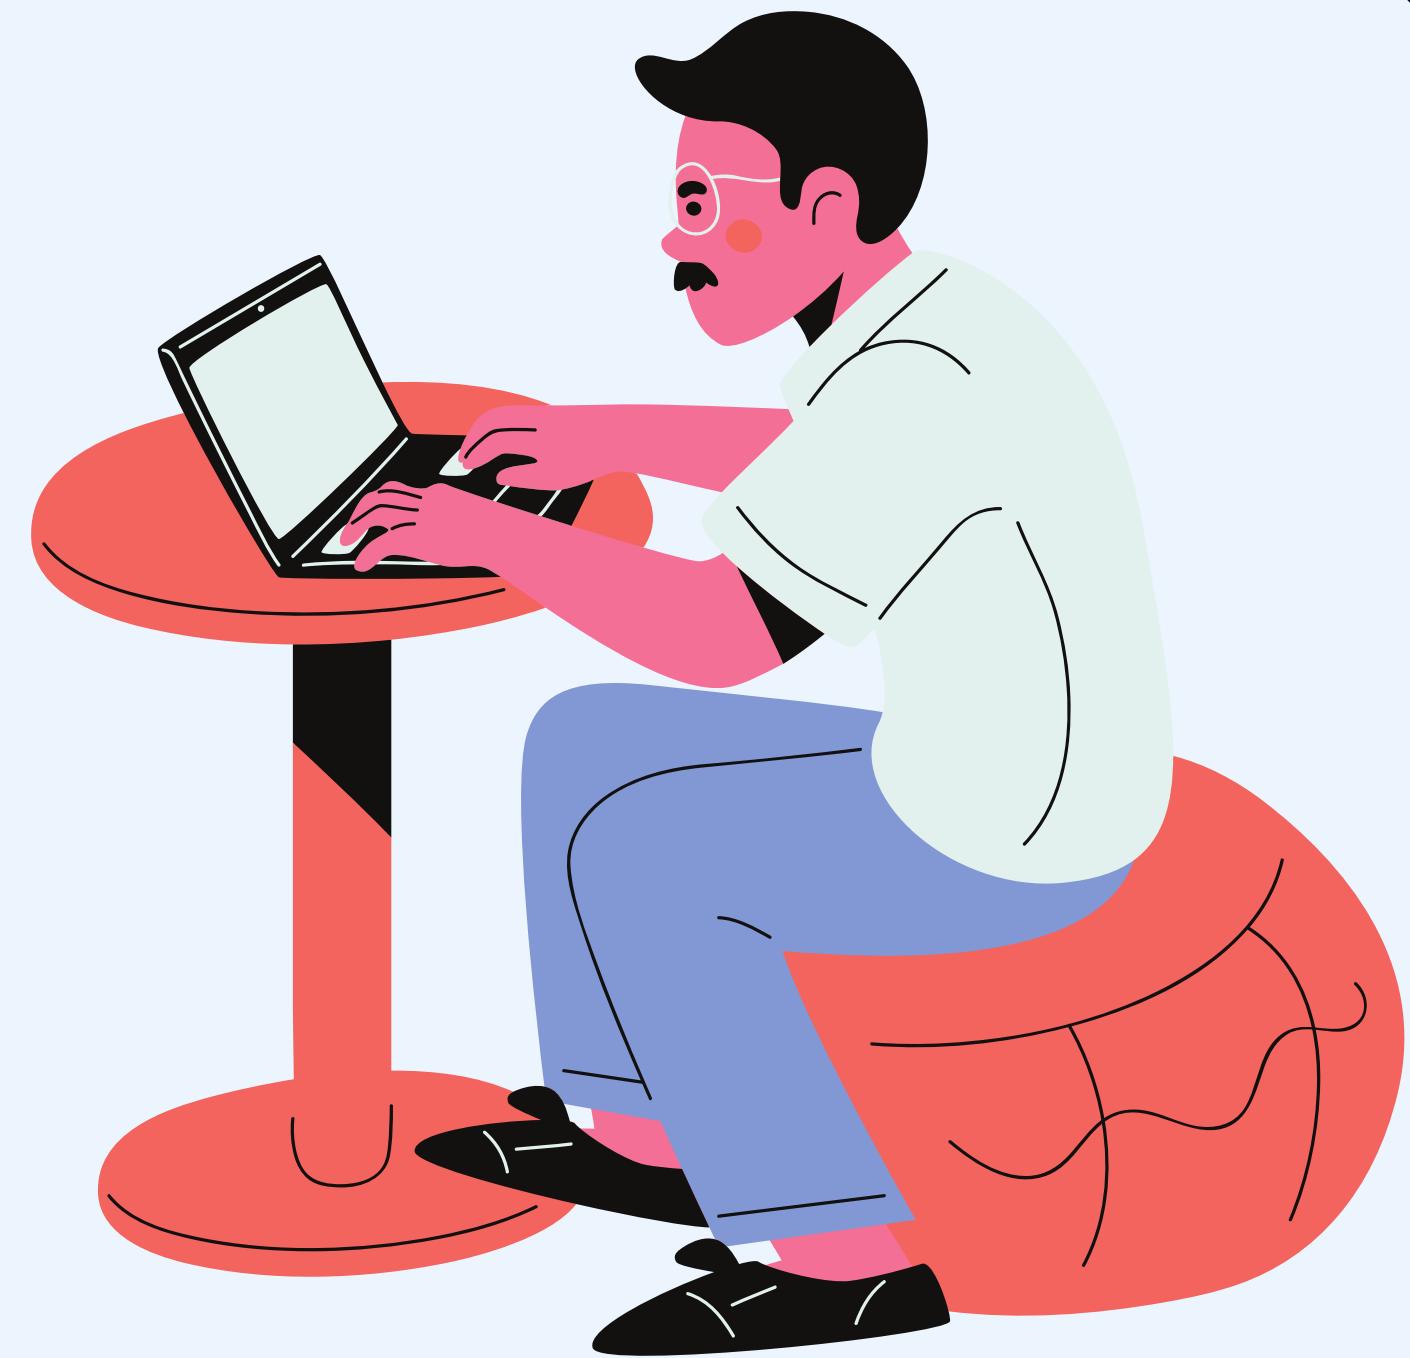


## Summary

- Mongoose custom validation: Flexible, beyond built-in validators
- Enforce unique, complex rules, or involve external APIs
- Custom validation for all your specific needs



# DEMO TIME





# ASYNC CUSTOM VALIDATORS



mongoose

MONGOOSE VALIDATION



# ASYNCHRONOUS CUSTOM VALIDATORS



## High-Level Explanation

- Asynchronous custom validators in Mongoose: Perform async operations
- Use `validate` with Promise for asynchronous logic
- Ideal for database queries, API calls, async validation



## Analogue

- Analogy: Asynchronous validation as party entry
- Bouncer checks list online (async)
- Confirmation depends on list loading



## Best Practices

- Handle Promise rejections to prevent unhandled errors.
- Keep operations lightweight to maintain application speed.
- Implement timeouts to prevent indefinite hangs in async operations.



## Deep Dive

- Asynchronous custom validator returns a Promise.
- Resolve with `true` for success, `false` for failure.
- Use `async/await` for async operations.
- Validation failure results in rejection and error thrown.



## When to use?

- External database checks for unique values are required.
- Validation relies on external API responses.
- I/O operations are necessary for your validation logic.

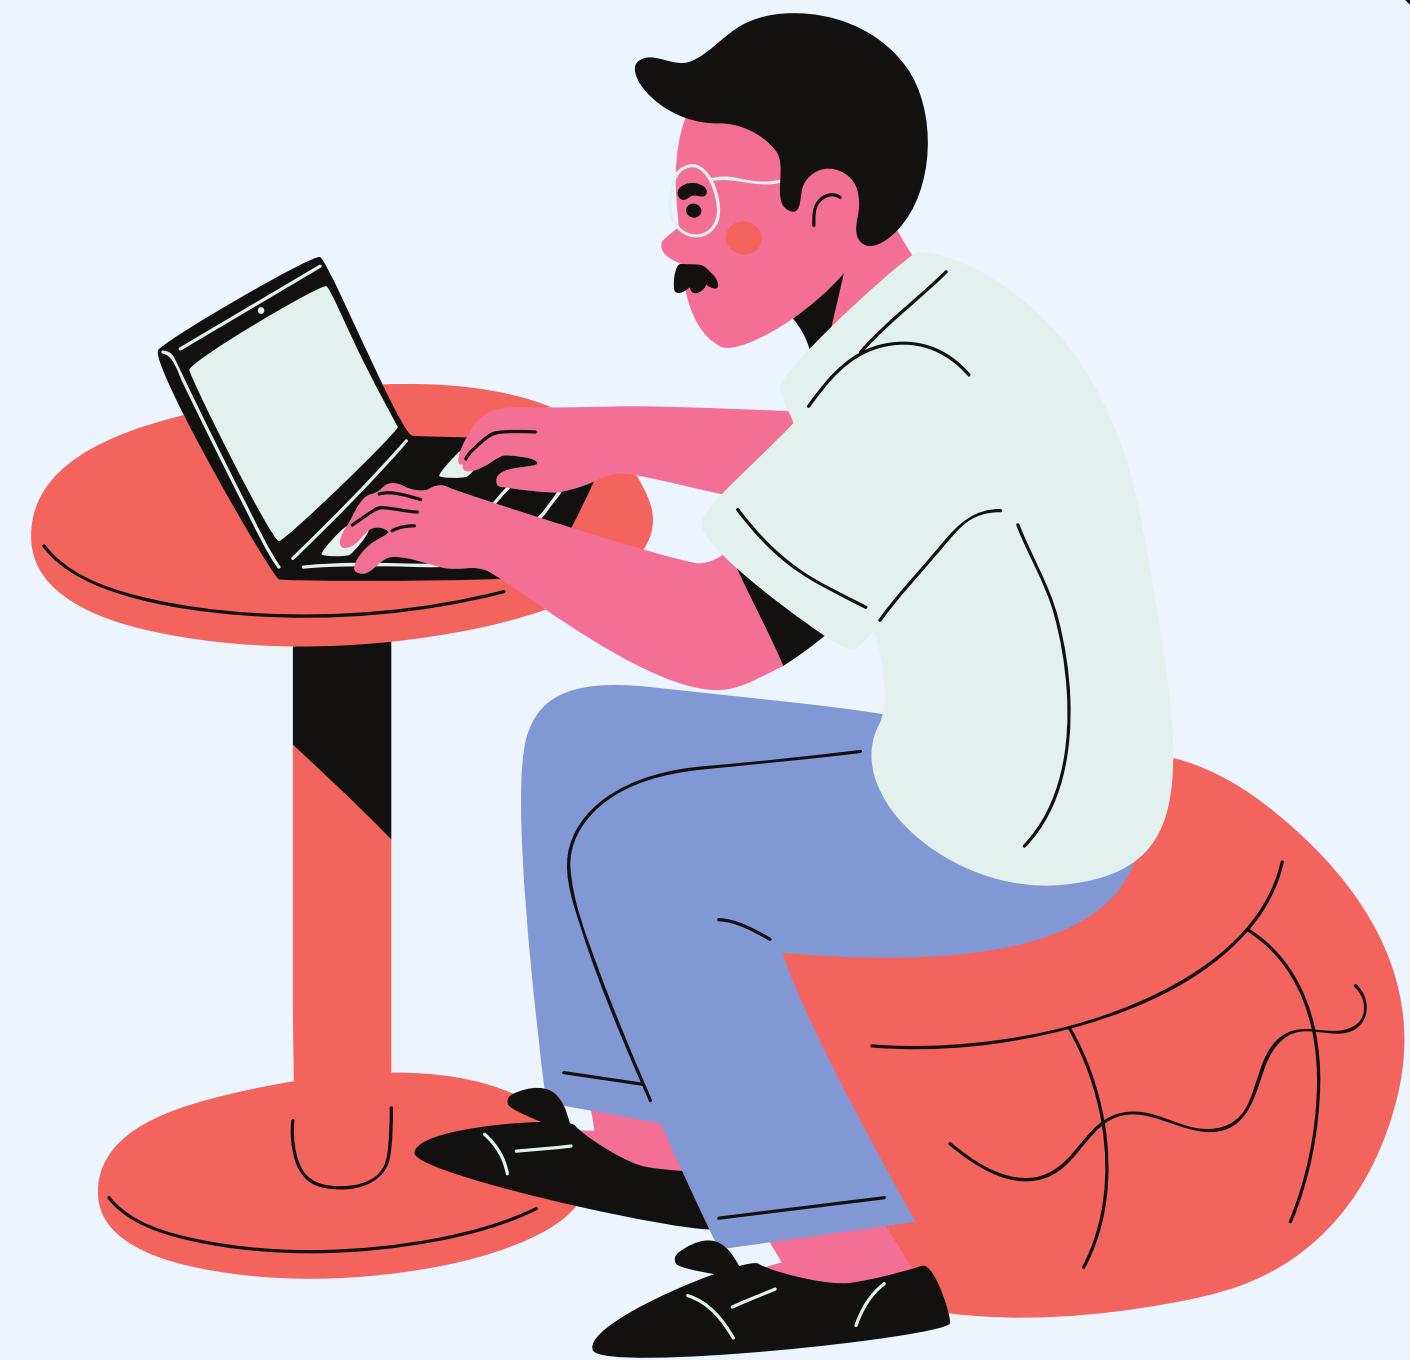


## Summary

- Asynchronous custom validators: Useful for external/time-consuming logic
- Implemented with Promise in `validator` or `async/await`



# DEMO TIME





# DATA SANITIZATION TECHNIQUES



mongoose

MONGOOSE VALIDATION

# DATA SANITIZATION TECHNIQUES



## High-Level Explanation

- Data sanitization: Clean, validate data before processing/saving
- Enforces quality standards, prevents malicious code



## Analogue

- Analogy: Database as a refrigerator
- Data sanitization as checking and cleaning items
- Ensure only safe, quality data is stored



## Best Practices

- Never trust user input; sanitize before processing.
- Utilize trusted libraries for common sanitization.
- Validate data types, length, format, and range.
- Use parameterized SQL queries to thwart SQL injection.
- Apply the "Principle of Least Privilege" to limit accepted data.



## Deep Dive

- Data sanitization layers: Client to server-side
- Techniques: Remove special chars, type conversion, allow-lists
- Prevents attacks like SQL injection, XSS, injection
- Mongoose, ExpressJS: Middleware, custom validators, libraries



## When to use?

- When handling external or user-generated data.
- For sensitive data with strict formatting (e.g., emails, URLs).
- To protect against malicious activities in database transactions.



## Summary

- Data sanitization: Clean, validate data for quality, security.
- Essential for application security, data integrity.



# DEMO TIME



# VALIDATOR LIBRARY



mongoose

MONGOOSE VALIDATION



# THE `SET` & `GET` SCHEMATYPE OPTIONS



mongoose

MONGOOSE VALIDATION



# THE `SET` AND `GET` SCHEMATYPE OPTIONS



## High-Level Explanation

- `set` option: Transforms data before saving to the database.
- `get` option: Transforms data after retrieval but before use.



## Analogue

- Analogy: Digital locker with `set` and `get`
- `set`: Locker organizes items when you put them in.
- `get`: Locker gift-wraps items when you take them out.



## Best Practices

- Setters: Avoid complex validation; focus on transformation.
- Note: Setters run each time property is set, not just on save.
- Getters: Avoid sensitive logic exposed to clients.
- Use sparingly to prevent unexpected behavior, debugging issues.



## Deep Dive

- Setters (`set`): Functions called when setting property values.
- Standardize data formatting (e.g., lowercase, trim).
- Getters (`get`): Functions after retrieving data but before use.
- Modify values (e.g., append domain to username).



## When to use?

- Setters: Sanitize, format data before saving to DB.
- Getters: Format, combine fields for app use without altering DB.

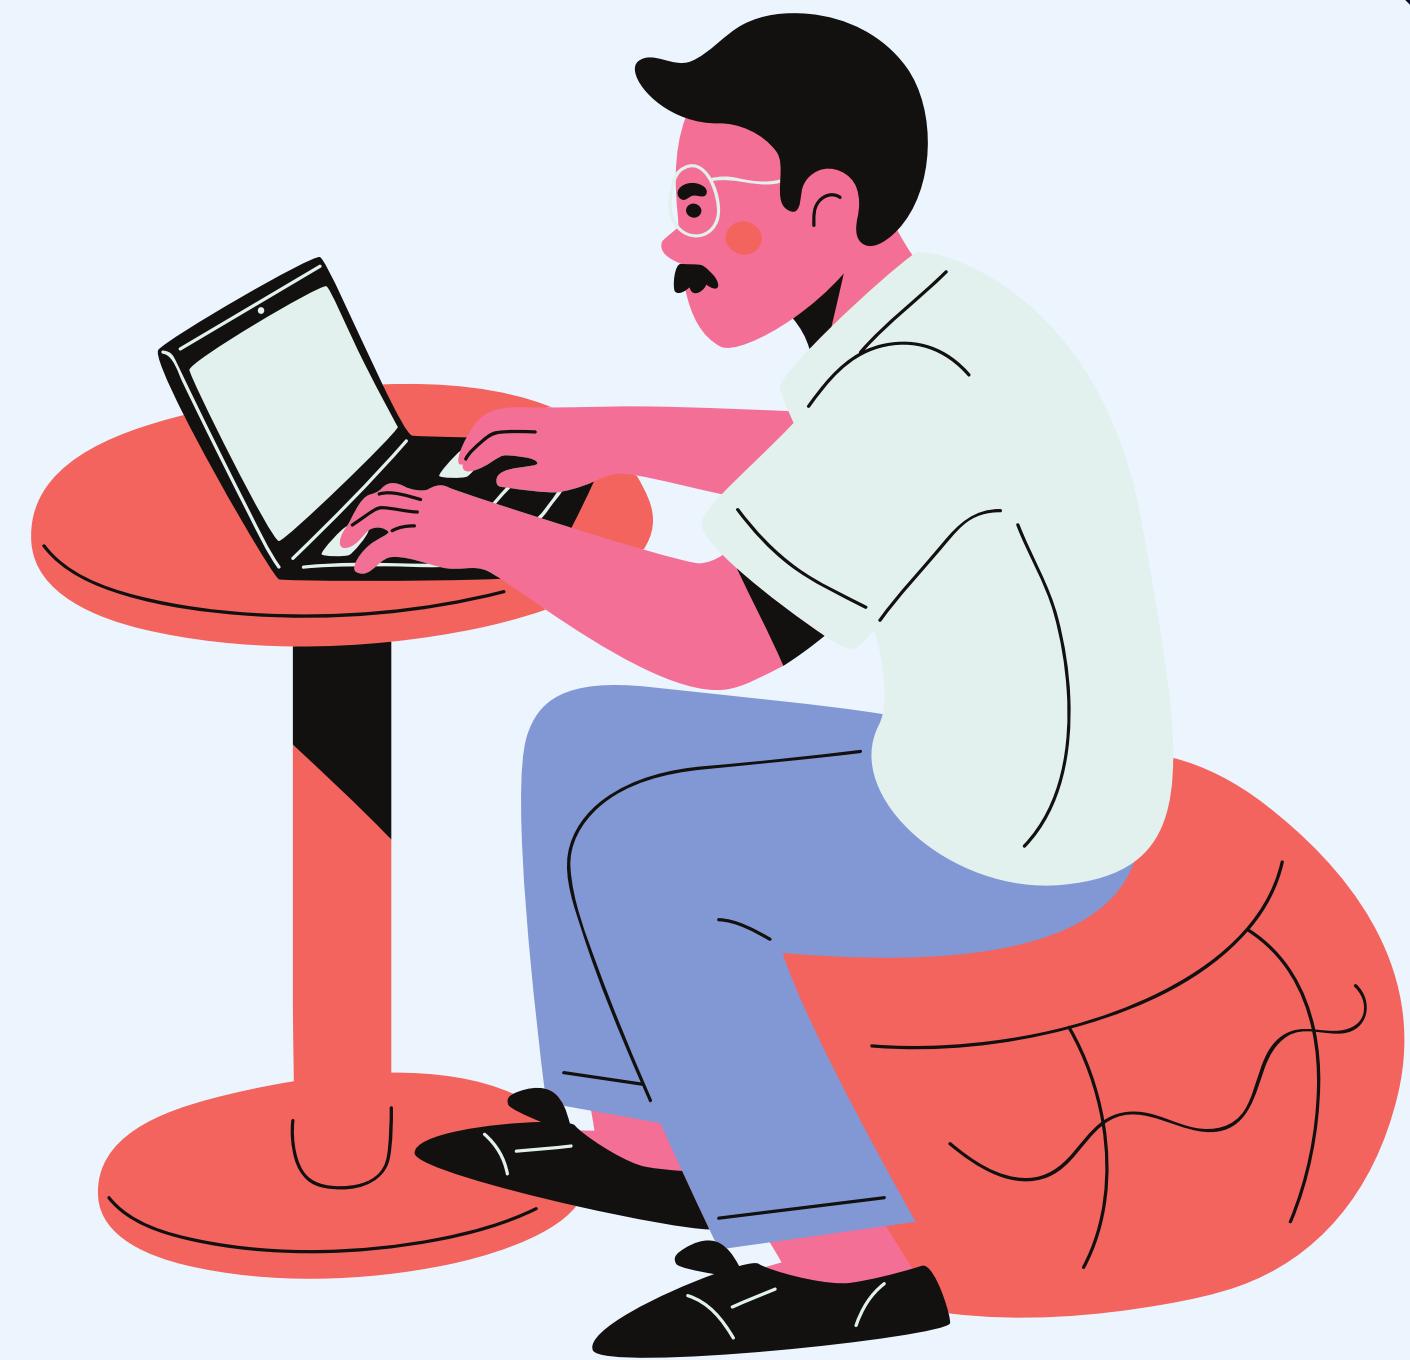


## Summary

- `set` and `get` in Mongoose schemas: Data transformation tools
- Setters prepare, sanitize before storage
- Getters polish, adjust data during retrieval



# DEMO TIME



# DATA MODELLING

**Overview**

mongoose



# DATA MODELLING OVERVIEW



## High-Level Explanation

- Data modeling: Defines data shape, relationships, integrity.
- Mongoose schemas: Specify MongoDB document structure.



## Deep Dive

- Data modeling: High-level design, relationships, integrity, DB-agnostic.
- Mongoose schemas: Specific for MongoDB, define types, validation, structure.



## Analogue

- Analogy: Data modeling as house furniture arrangement.
- Schema as furniture instruction manual.
- Defines parts and assembly for specific furniture.



## When to use?

- Data modeling: Initial app design, high-level data plan.
- Schemas: Implement high-level design in MongoDB with Mongoose.



## Best Practices

- Data modeling: Grasp business requirements thoroughly.
- Schemas: Define types, set up validation for data integrity.



## Summary

- Data modeling: Design data structure and relations.
- Schemas in Mongoose: Implement design for MongoDB documents.



# EMBEDDED DOCUMENTS



mongoose

DATA MODELLING

# UNDERSTANDING EMBEDDED DOCUMENTS



## High-Level Explanation

- Embedded documents: Sub-docs within a parent doc.
- Suitable for nested structures, e.g., address in user profile.



## Analogue

- Analogy: Embedded docs as backpack pockets.
- Backpack = Main item (parent doc).
- Pockets = Embedded docs holding specific items.
- Pockets are part of backpack, not separate.



## Best Practices

- Monitor document size limit.
- Avoid embedded docs for frequently independent queries.
- Acknowledge complexity in updates to embedded docs within parent.



## Deep Dive

- Define embedded documents in Mongoose with separate schema.
- Include them in parent schema as array or nested object.
- Embedded docs have fields, validation, methods.
- No separate database collection; only in parent doc.
- Efficient reads but potential document size limits.



## When to use?

- Use embedded docs when:
  - Clear hierarchical data.
  - Frequent retrieval of parent and embedded data.
  - Embedded data within MongoDB's size limit (16MB).



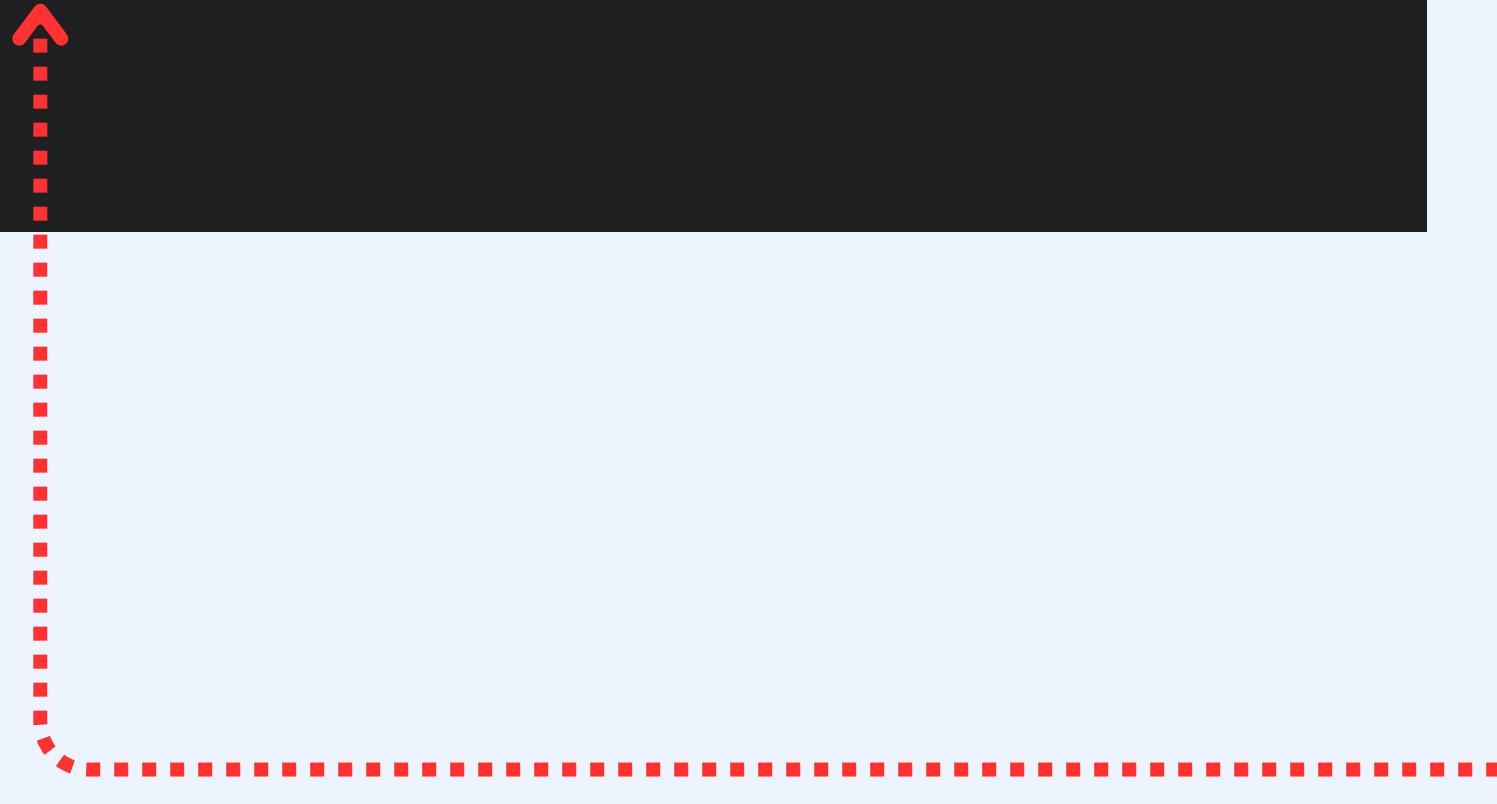
## Summary

- Embedded docs in Mongoose: Sub-docs in parent doc.
- Ideal for nested, hierarchical data.
- Limitations include MongoDB's doc size limit.



## parent document schema

```
const UserSchema = new Schema({  
  name: String,  
  email: String,  
  address: AddressSchema // Embedding AddressSchema  
});
```



## Embedded document schema

```
const AddressSchema = new Schema({  
  street: String,  
  city: String,  
  state: String,  
  zip: Number  
});
```



# DEMO TIME





# ARRAYS OF SUB-DOCUMENTS



mongoose

DATA MODELLING

# WORKING WITH ARRAYS OF SUBDOCUMENTS



## High-Level Explanation

- Arrays of subdocuments in Mongoose manage embedded docs.
- Useful for one-to-many relationships within one parent doc.
- Manipulate arrays for efficient data management (add, update, delete subdocs).



## Analogue

- Analogy: Embedded docs as backpack pockets.
- Backpack = Main item (parent doc).
- Pockets = Embedded docs holding specific items.
- Pockets are part of backpack, not separate.



## Best Practices

- Watch for MongoDB's 16MB doc size limit.
- Large subdoc arrays: Consider pagination or separate collections.
- Use Mongoose validation for subdocs to ensure data integrity.



## Deep Dive

- Create arrays of subdocs with `[]` notation in schema.
- Add subdocs using `push` or by assigning a new array.
- Update by finding parent and using array methods or Mongoose subdoc API.
- Delete by retrieving parent and using `splice` or Mongoose methods.
- Query subdoc fields, but performance may differ from standalone docs.



## When to use?

- Use arrays of subdocs for:
  - Strongly related nested data.
  - Collections within doc size limits.
  - Operations often need entire sub-array, reducing complex queries.



## Summary

- Arrays of subdocs in Mongoose: Manage one-to-many within one doc.
- Great for smaller, closely related data sets.
- Challenges: Complex querying, 16MB doc size limit.
- Awareness of constraints informs design decisions.

# WORKING WITH ARRAYS OF SUBDOCUMENTS



## Sub document schema

```
const SubdocumentSchema = new mongoose.Schema({  
  field1: String,  
  field2: Number  
});
```



## Parent schema

```
const ParentSchema = new mongoose.Schema({  
  subdocuments: [SubdocumentSchema]  
});
```



# DEMO TIME



# LIMITATIONS OF EMBEDDED DOCUMENTS

mongoose



DATA MODELLING



# LIMITATIONS OF EMBEDDED DOCUMENTS



## High-Level Explanation

- Embedded docs in Mongoose: Store nested data in MongoDB doc.
- Manage relationships without joins.
- Limitations: Doc size constraints, query complexity, limited ops support.
- Consider pros and cons for design decisions.



## Deep Dive

- MongoDB's 16MB BSON document size limit applies to embedded docs.
- Query complexity increases with nested fields.
- Some operations (e.g., "distinct") less straightforward.
- Redundancy and data integrity challenges.
- ACID transactions less performant on deeply nested embedded docs.



## Analogue

- Analogy: Computer's game folders.
- Main folder = "My Games."
- Sub-folders = Embedded docs.
- 16MB limit = Computer's "no more space."
- Complexity = Treasure hunt for specific info in embedded docs.

# DATA REFERENCING

mongoose



DATA MODELLING



# UNDERSTANDING DATA REFERENCING



## High-Level Explanation

- Data referencing in Mongoose: Linking docs from different collections.
- Used for many-to-many, one-to-many relationships.
- Avoids size limits, supports complex queries.



## Analogue

- Analogy: Toy box (collection) and friends (another collection).
- Referencing: Friend's name on toy list, not on toy (no embedding).
- Query: List helps find who has which toy.



## Best Practices

- `populate()`: Easy data retrieval but can be slow for large datasets.
- For frequent multi-collection references, optimize data architecture (indexing).
- Beware "callback hell" with nested queries for multiple references.



## Deep Dive

- Data referencing in Mongoose: Two types - manual and population.
- **Manual referencing:** Store related `\\_id`, manual queries.
- **Population:** `populate()` auto-replaces `\\_id`, simpler retrieval.
- **Normalization:** Data separate in collection, not duplicated.
- Contrast: Denormalized data has subdocs embedded in parent.



## When to use?

- Use data referencing when:
  - Referenced data is large (exceeds 16MB limit).
  - Data is shared among multiple docs.
  - Complex queries on referenced data needed.



## Summary

- Data referencing in Mongoose: Relate separate docs.
- Manual referencing = query control.
- `populate()`: Simplifies data retrieval.
- Strategic use enhances database efficiency.



# UNDERSTANDING DATA REFERENCING

## Authors document schema

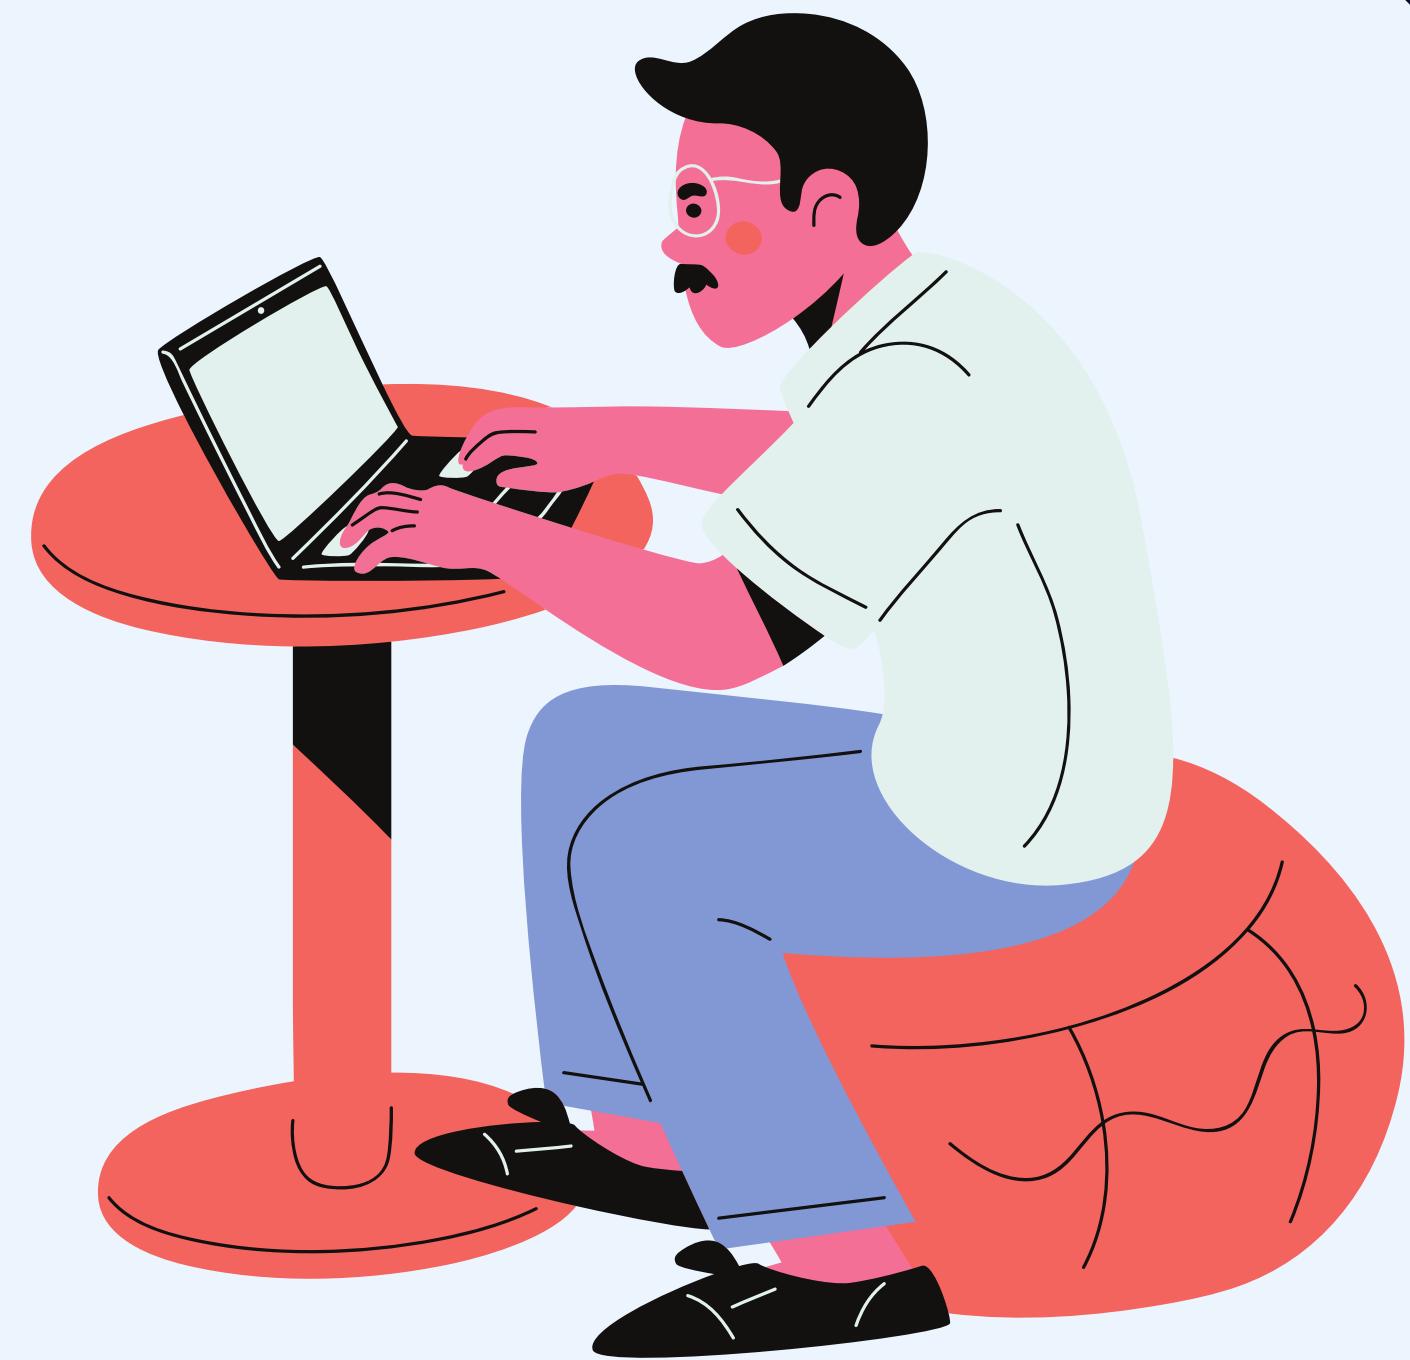
```
const AuthorSchema = new mongoose.Schema({  
  name: String,  
  birthYear: Number  
});
```

## Books document schema

```
const BookSchema = new mongoose.Schema({  
  title: String,  
  author: {  
    type: mongoose.Schema.Types.ObjectId,  
    ref: 'Author'  
  }  
});
```



# DEMO TIME



# THE `REF` & POPULATE PROPERTIES

mongoose



DATA MODELLING



# THE `REF` PROPERTY AND POPULATION



## High-Level Explanation

- `ref` specifies the referenced model in a schema.
- `populate` replaces Object IDs with data from referenced model.
- Similar to SQL JOIN, simplifying data retrieval.



## Deep Dive

- `ref` property: Used to specify the referenced model in a schema.
- `populate` method: Replaces Object IDs with data from the referenced model.
- Query-level population: Allows for dynamic data retrieval as needed.



## Analogue

- The `ref` property specifies the referenced model for a field.
- `populate()` fetches data from the referenced model, akin to opening a treasure chest to reveal its contents in a game.



## When to use?

- Use `ref` in schema to specify referenced model.
- `populate` dynamically replaces Object IDs with data.
- Useful for normalized data, complex updates, and on-demand loading.



## Best Practices

- Use population sparingly for large datasets.
- Selective field population can limit populated fields.

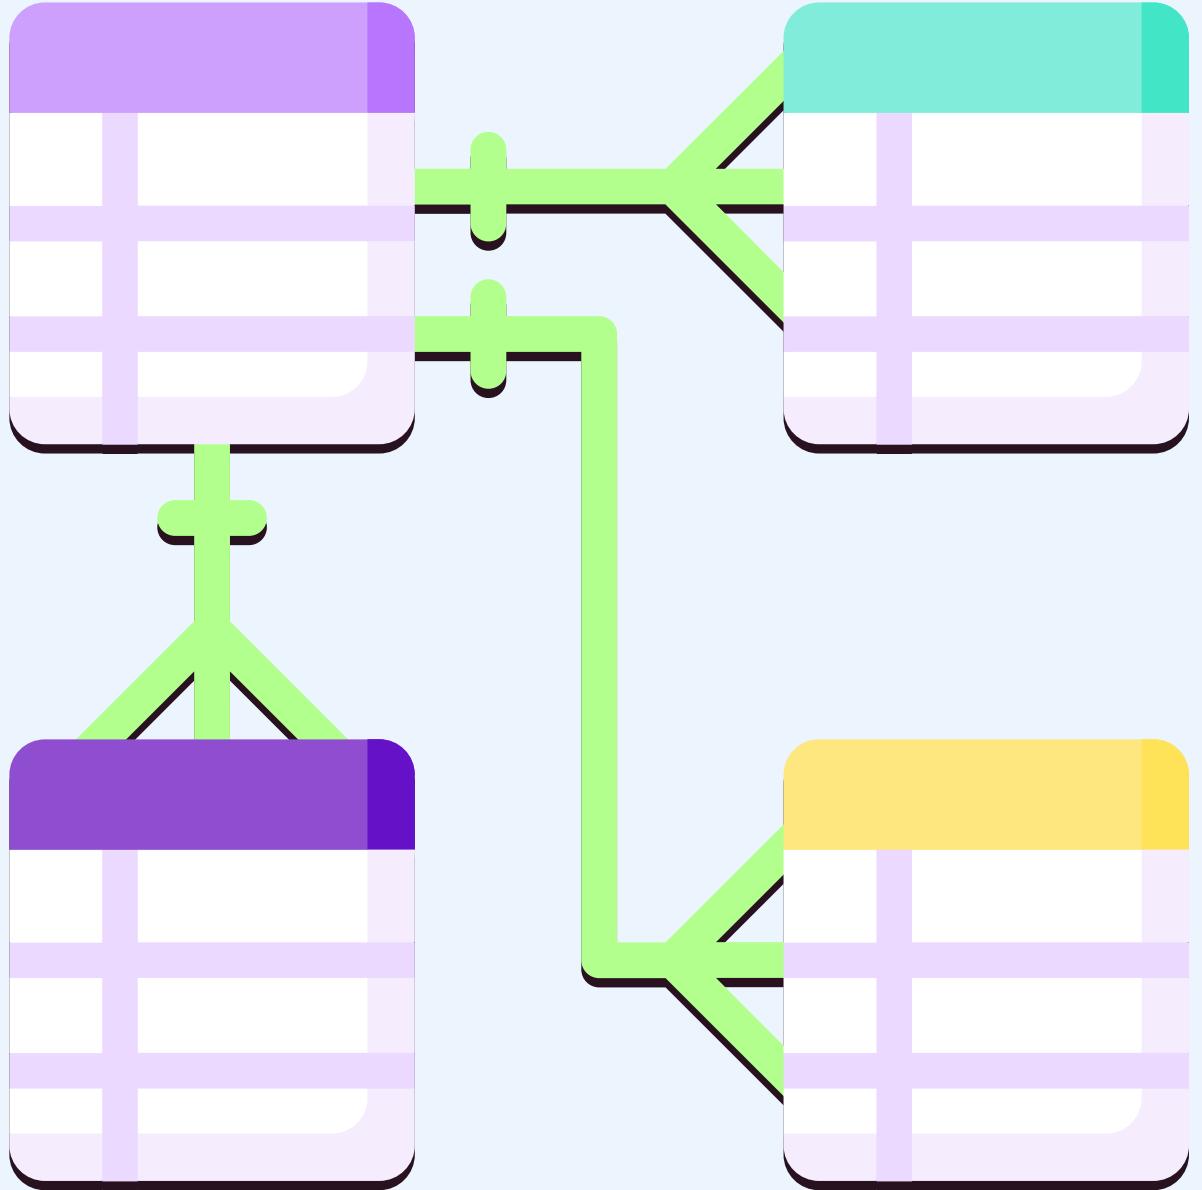


## Summary

- `ref` property specifies the referenced model for an Object ID field.
- `populate` uses this reference to replace IDs with actual docs.
- Essential for managing complex data relationships in Mongoose.



# DATA RELATIONSHIP



mongoose

MONGOOSE DATA RELATIONSHIP



# ONE-ONE RELATIONSHIPS



mongoose

MONGOOSE DATA RELATIONSHIP



# ONE-TO-ONE RELATIONSHIPS



## High-Level Explanation

- One-to-one relationship: Single doc in one collection matches one in another.
- Implement via Embedded Docs or References in Mongoose.



## Deep Dive

- Embedded: Store related doc as sub-doc within parent.
- Reference-based: Use ObjectId reference in another collection.
- Choice depends on app requirements: access patterns, change rate, data size.



## Analogue

- Embedded: "About Me" content directly on profile page.
- Reference-based: "About Me" content on a separate page, linked from the profile.



## When to use?

- Frequent access to associated data together.
- True one-to-one connection between documents.
- Data size won't exceed MongoDB's document size limit.



## Best Practices

- Embed data for read-heavy and stable data.
- Use references for write-heavy or dynamic data.
- Base decision on application's data access patterns.



## Summary

- One-to-one relationships in MongoDB via embedding or referencing.
- Decision hinges on app needs, like data access and change frequency.



# ONE-TO-ONE RELATIONSHIPS

## Embedded Approach

```
Users Collection
  └ User Document
    └ _id: ObjectId("...")  

    └ name: "John Doe"  

    └ email: "john.doe@example.com"  

    └ address: Embedded Document
      └ street: "123 Main St"  

      └ city: "Springfield"  

      └ zipCode: "62704"
```



# ONE-TO-ONE RELATIONSHIPS

## Reference-Based Approach

### Users Collection

#### └ User Document

```
    └ _id: ObjectId("...")
```

```
    └ name: "John Doe"
```

```
    └ email: "john.doe@example.com"
```

```
    └ address: ObjectId("...")
```

### Addresses Collection

#### └ Address Document

```
    └ _id: ObjectId("...")
```

```
    └ street: "123 Main St"
```

```
    └ city: "Springfield"
```

```
    └ zipCode: "62704"
```



# DEMO TIME



# ONE-MANY MANY-ONE RELATIONSHIPS

mongoose



MONGOOSE DATA RELATIONSHIP

# ONE-TO-MANY AND MANY-TO-ONE RELATIONSHIPS



## High-Level Explanation

- One-to-many: One document in Collection A associated with multiple in Collection B.
- Many-to-one: Multiple in Collection A refer to one in Collection B.
- Perspective shift differentiates the two.



## Analogue

- Classroom analogy: One-to-many – one teacher, many students.
- Student's perspective: Many students, one teacher – many-to-one.



## Best Practices

- Use arrays for one-to-many but watch document size limit.
- Avoid deep nested arrays for simpler queries/updates.
- In many-to-one, maintain referential integrity by verifying existence of referred-to document.



## Deep Dive

- **One-to-many:** Use arrays with ObjectIds in source collection.
  - E.g., Blog post with an array of comment references.
- **Many-to-one:** Multiple records refer to one in target collection.
  - E.g., Many comments belong to a single blog post, specified by a field.



## When to use?

- **One-to-many:** Good for logical grouping, infrequent access.
  - Data consistency less critical (no MongoDB transactions).
- **Many-to-one:** Useful when frequent access to a single document.
  - Logical to link multiple records to one central record.



## Summary

- One-to-many and many-to-one are two perspectives of the same relationship.
- One-to-many: Array of references in 'one' to 'many.'
- Many-to-one: 'Many' pointing back to the 'one.'



## One-to-Many Relationship

BlogPosts Collection

  └ BlogPost Document

    ├ \_id: ObjectId("...")

    ├ title: "My First Blog"

    └ comments: Array of ObjectIds

Comments Collection

  └ Comment Document 1

    ├ \_id: ObjectId("...")

    ├ text: "Great post!"

    └ blogPostId: ObjectId("...")

  └ Comment Document 2

    ├ \_id: ObjectId("...")

    ├ text: "Thanks for sharing!"

    └ blogPostId: ObjectId("...")



## Many-to-One Relationship

Comments Collection

- └─ Comment Document 1
  - ├─ \_id: ObjectId("...")
  - ├─ text: "Great post!"
  - └─ blogPostId: ObjectId("...") <--
- └─ Comment Document 2
  - ├─ \_id: ObjectId("...")
  - ├─ text: "Thanks for sharing!"
  - └─ blogPostId: ObjectId("...")

BlogPosts Collection

- └─ BlogPost Document
  - ├─ \_id: ObjectId("...")
  - ├─ title: "My First Blog"
  - └─ comments: Array of ObjectIds



# DEMO TIME





# MANY TO MANY RELATIONSHIPS



mongoose

MONGOOSE DATA RELATIONSHIP



# MANY-TO-MANY RELATIONSHIPS

## Many-to-Many Relationship

### Students Collection

```
└ Student Document 1
    └ _id: ObjectId("...")  

    └ name: "Alice"  

    └ enrolledCourses: Array of ObjectIds ---> enrolledStudents: Array of Objects  

        └ ObjectId("...")  

        └ ObjectId("...")  

└ Student Document 2
    └ _id: ObjectId("...")  

    └ name: "Bob"  

    └ enrolledCourses: Array of ObjectIds ---> enrolledStudents: Array of Objects  

        └ ObjectId("...")  

        └ ObjectId("...")
```

### Courses Collection

```
└ Course Document 1
    └ _id: ObjectId("...")  

    └ name: "Math 101"  

    └ ObjectId("...")  

    └ ObjectId("...")  

└ Course Document 2
    └ _id: ObjectId("...")  

    └ name: "History 101"  

    └ ObjectId("...")  

    └ ObjectId("...")
```

# MANY-TO-MANY RELATIONSHIPS



## High-Level Explanation

- Many-to-many relationship: Multiple docs in one collection associate with multiple docs in another.
- Complex compared to one-to-many and many-to-one relations.



## Analogue

- Social media platforms like Instagram: Classic many-to-many relationship.
- You follow multiple people, and multiple people follow you.
- Mutual, bidirectional relationship.



## Best Practices

- Watch array size due to MongoDB's 16MB limit.
- Complex relationships? Consider a join collection.
- Maintain data consistency and integrity during updates/deletes.



## Deep Dive

- Many-to-many relationship in Mongoose: Arrays of ObjectIds in both collections.
- Example: Students with course ObjectIds, Courses with student ObjectIds.
- Allows multiple students in multiple courses.



## When to use?

- Use many-to-many relationships when:
  - Complex associations between document sets.
  - Non-hierarchical, non-singular relationships.
  - Frequent need for querying these relationships.

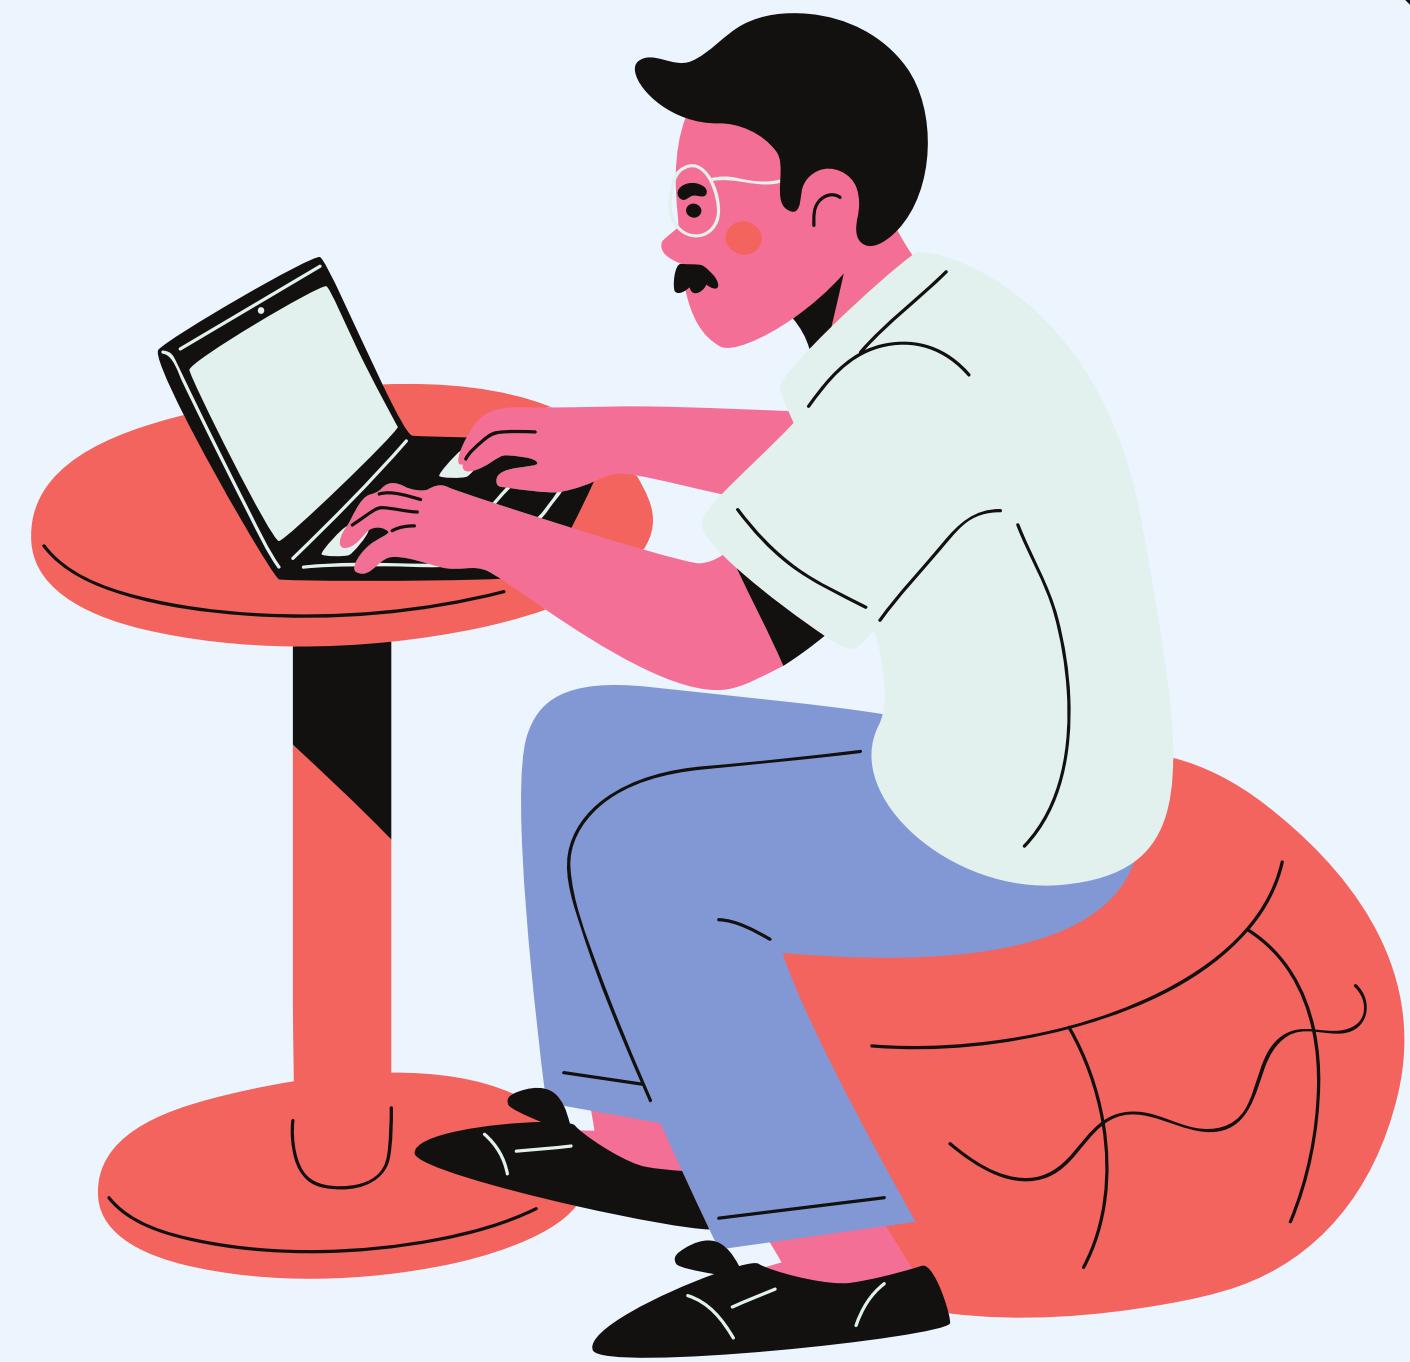


## Summary

- Many-to-many relationships: Multiple docs related to multiple docs.
- Implement with arrays of references or separate join collection.



# DEMO TIME



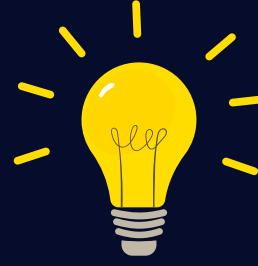


# DATA AGGREGATION

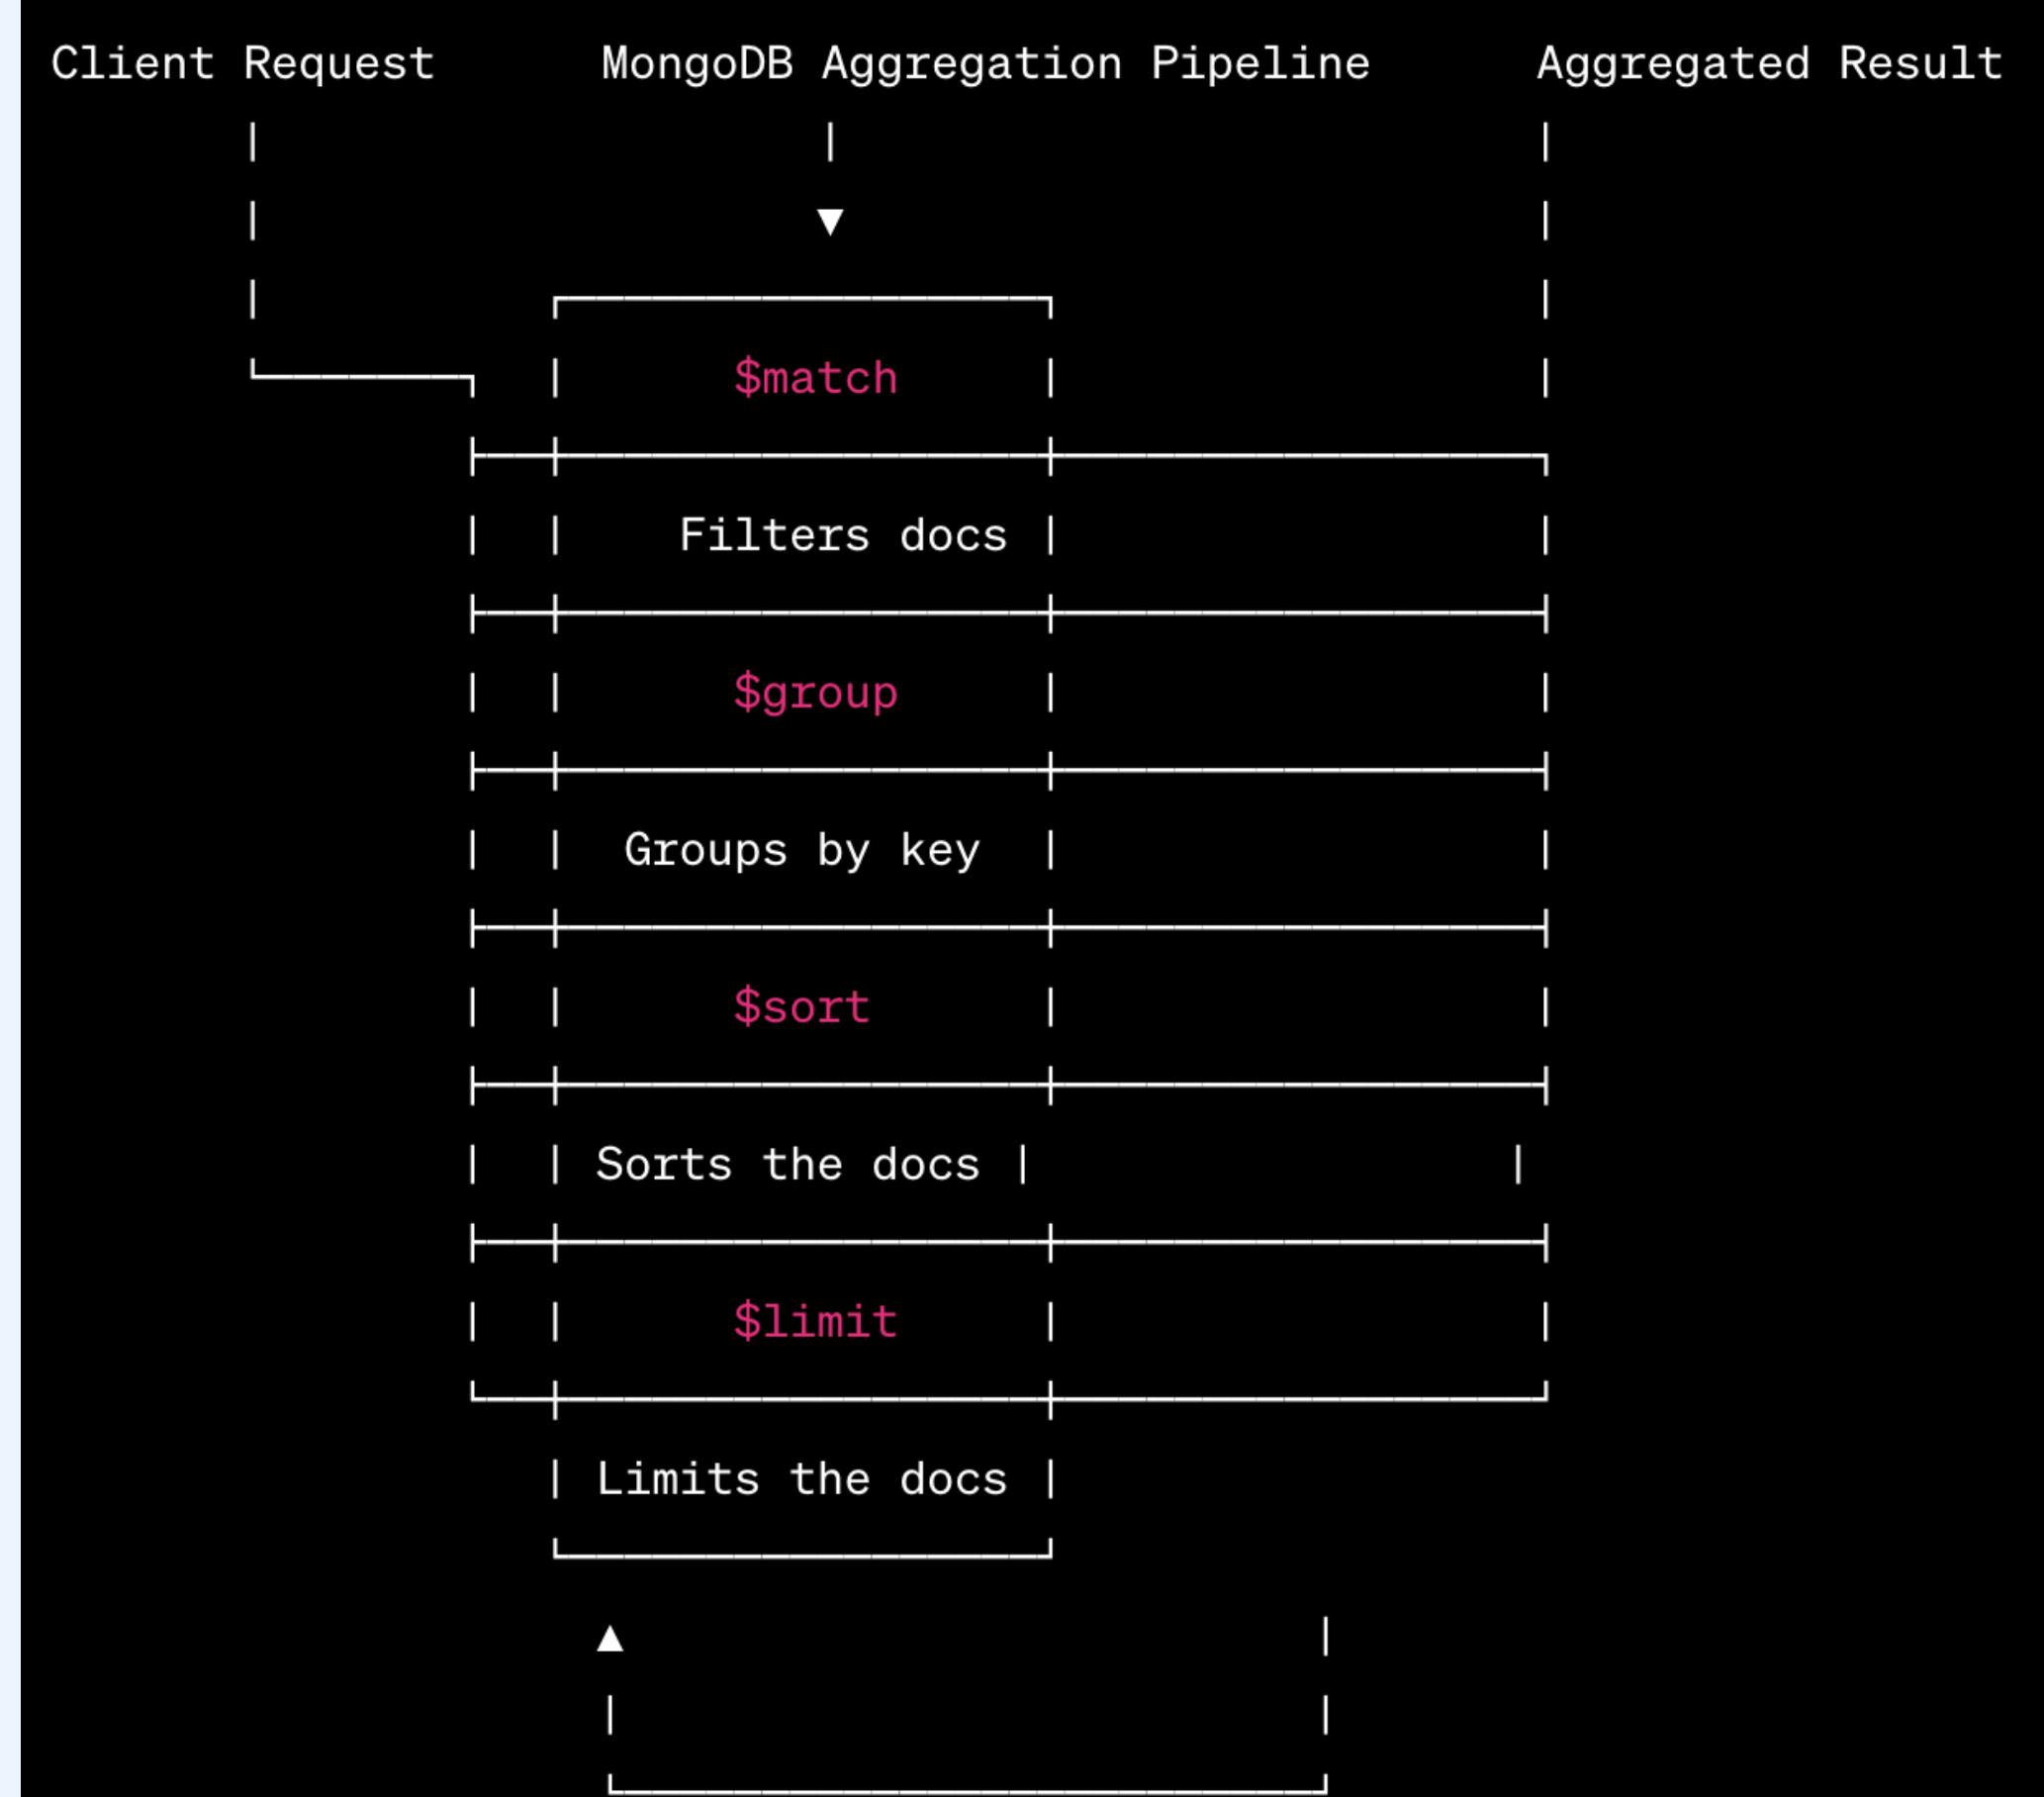


mongoose

MONGOOSE DATA RELATIONSHIP



# DATA AGGREGATION





# DATA AGGREGATION



## High-Level Explanation

- Data aggregation in MongoDB: Transform data via operations.
- Compute derived data, stats, manipulations.
- Original dataset remains unchanged.



## Analogue

- Analogy: Creating a playlist from music collection.
- Aggregation stages:
  1. `'\$match`': Pick rock songs.
  2. `'\$group`': Categorize by artist.
  3. `'\$sort`': Sort by song length.
  4. `'\$limit`': Select top 10 songs.



## Best Practices

- Optimize with indexes for `'\$match`' and `'\$sort`'.
- Control document count per stage for memory.
- Monitor performance for resource-intensive aggregations.



## Deep Dive

- MongoDB aggregation pipeline: Complex data transformation.
- Ordered stages process documents to aggregated results.
- Common stages: `'\$match`', `'\$group`', `'\$sort`', `'\$limit`', etc.
- Mongoose: Use `'.aggregate()`` method on a model.



## When to use?

- Use MongoDB aggregation when you:
  - Summarize data for analytics.
  - Transform documents into different forms.
  - Compute derived data or statistics.



## Summary

- Data aggregation in MongoDB/Mongoose: Powerful for data manipulation.
- Aggregation pipeline: Sequence of ops for filtering, grouping, sorting, and aggregating.
- Generates aggregated results from documents.



# DEMO TIME





# MONGOOSE MIDDLEWARES OVERVIEW

mongoose



MONGOOSE MIDDLEWARES



# MONGOOSE MIDDLEWARES OVERVIEW



## High-Level Explanation

- Mongoose middlewares: Functions in Mongoose document lifecycle.
- Hook into stages like validation, saving, finding.
- Enable custom logic, data modification.



## Analogue

- Analogy: Mongoose middlewares as video game checkpoints.
- Check, modify data before proceeding to next stage.
- Ensure conditions are met before advancing.



## Best Practices

- Keep them small and focused, performing a single task.
- Handle errors in asynchronous operations.
- Be cautious when modifying queries or data to avoid unintended consequences.



## Deep Dive

Mongoose middlewares are categorized into different types:

- **Document Middlewares:** Operate on individual documents (e.g., `save`, `validate`, `remove`).
- **Query Middlewares:** Operate on queries (e.g., `find`, `update`).
- **Aggregate Middlewares:** Work with aggregation pipelines.
- **Model Middlewares:** Apply to model's static functions.

They can be synchronous or asynchronous, registered using `pre` and `post` hooks.



## When to use?

- Document middlewares: Single document actions (validation, modification).
- Query middlewares: Query modification, logging.
- Aggregate middlewares: Complex data transformation.
- Model middlewares: Collection-wide actions (e.g., data seeding).



## Summary

- Mongoose middlewares: Insert custom logic at data operation stages.
- Apply at document, query, and other levels.
- Essential for effective data management in Mongoose apps.



# DOCUMENT MIDDLEWARES



mongoose

MONGOOSE MIDDLEWARES



# DOCUMENT MIDDLEWARES



## High-Level Explanation

- Mongoose Document middlewares (hooks): Custom logic pre/post events.
- Hooks: `init`, `validate`, `save`, `remove`.
- Trigger at various stages in document lifecycle.



## Analogue

- Analogy: Mongoose Document middlewares as theater crew.
- `init`: Preparing the stage before the show.
- `validate`: Ensuring everything is in place.
- `save`: Managing the action during the performance.
- `remove`: Cleaning up after the show.



## Best Practices

- Asynchronous: Employ `next()` for async continuation.
- Error Handling: Manage errors with try-catch or pass to `next()`.
- Immutable Fields: Don't alter immutable fields in `pre('save')` hooks.



## Deep Dive

Mongoose Document middlewares provide control and customization of operations:

- `init`: Triggers on document load.
- `validate`: Pre-validation.
- `save`: Pre/post save (before/after).
- `remove`: Pre/post remove (before/after deletion from collection).



## When to use?

- `init`: Set initial values or conditions.
- `validate`: Enforce complex validation rules.
- `save`: Automate tasks, manipulate data pre/post save.
- `remove`: Perform clean-up, delete associated records.



## Summary

- Mongoose Document middlewares: Custom logic at lifecycle stages
- Init, validate, save, remove hooks
- Enhance app robustness and maintainability

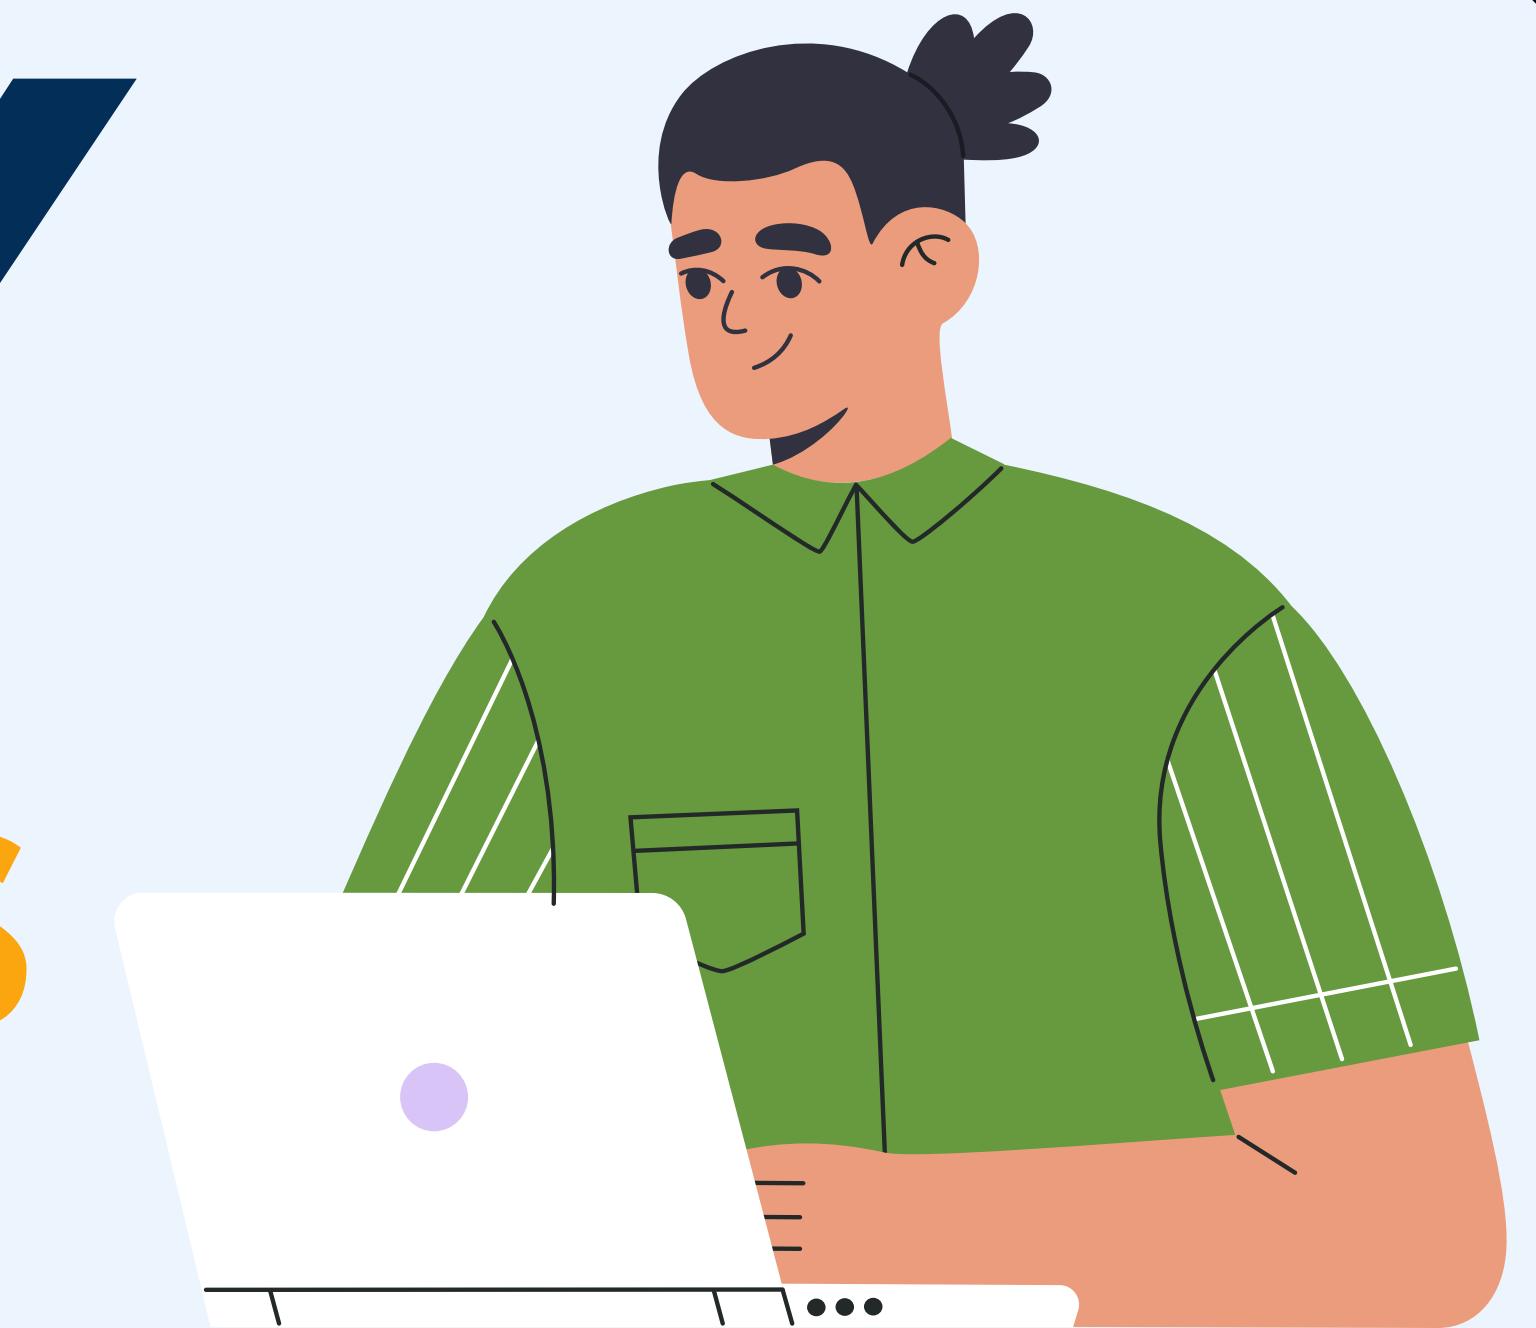


# DEMO TIME



# QUERY

## MIDDLEWARES



mongoose

MONGOOSE MIDDLEWARES



# QUERY MIDDLEWARES



## High-Level Explanation

- Mongoose Document middlewares (hooks): Custom logic pre/post events.
- Hooks: `init`, `validate`, `save`, `remove`.
- Trigger at various stages in document lifecycle.



## Analogue

- Analogy: Google Search and Mongoose query middlewares
- Pre-middleware: Filter out specific terms (like Google)
- Post-middleware: Count results for analytics (like Google)
- Mongoose middlewares for database operations



## Best Practices

1. Utilize `async/await` for async operations in middlewares.
2. Keep middlewares lightweight; avoid heavy computations for performance.
3. Carefully chain multiple middlewares to maintain code readability.
4. Always call `next()` in pre-middlewares to continue execution.



## Deep Dive

- Query middleware uses: Data normalization, validation, caching
- Operate on query objects, not documents
- Two types: "pre" (before) and "post" (after)
- Modify query with `pre('find')`, process results with `post('find')`



## When to use?

- Common use cases for Mongoose query middlewares:
1. Data Validation: Ensure queries meet specific criteria.
  2. Authorization: Check permissions before executing a query.
  3. Caching: Implement caching for query performance.
  4. Logging: Track database activities for auditing purposes.

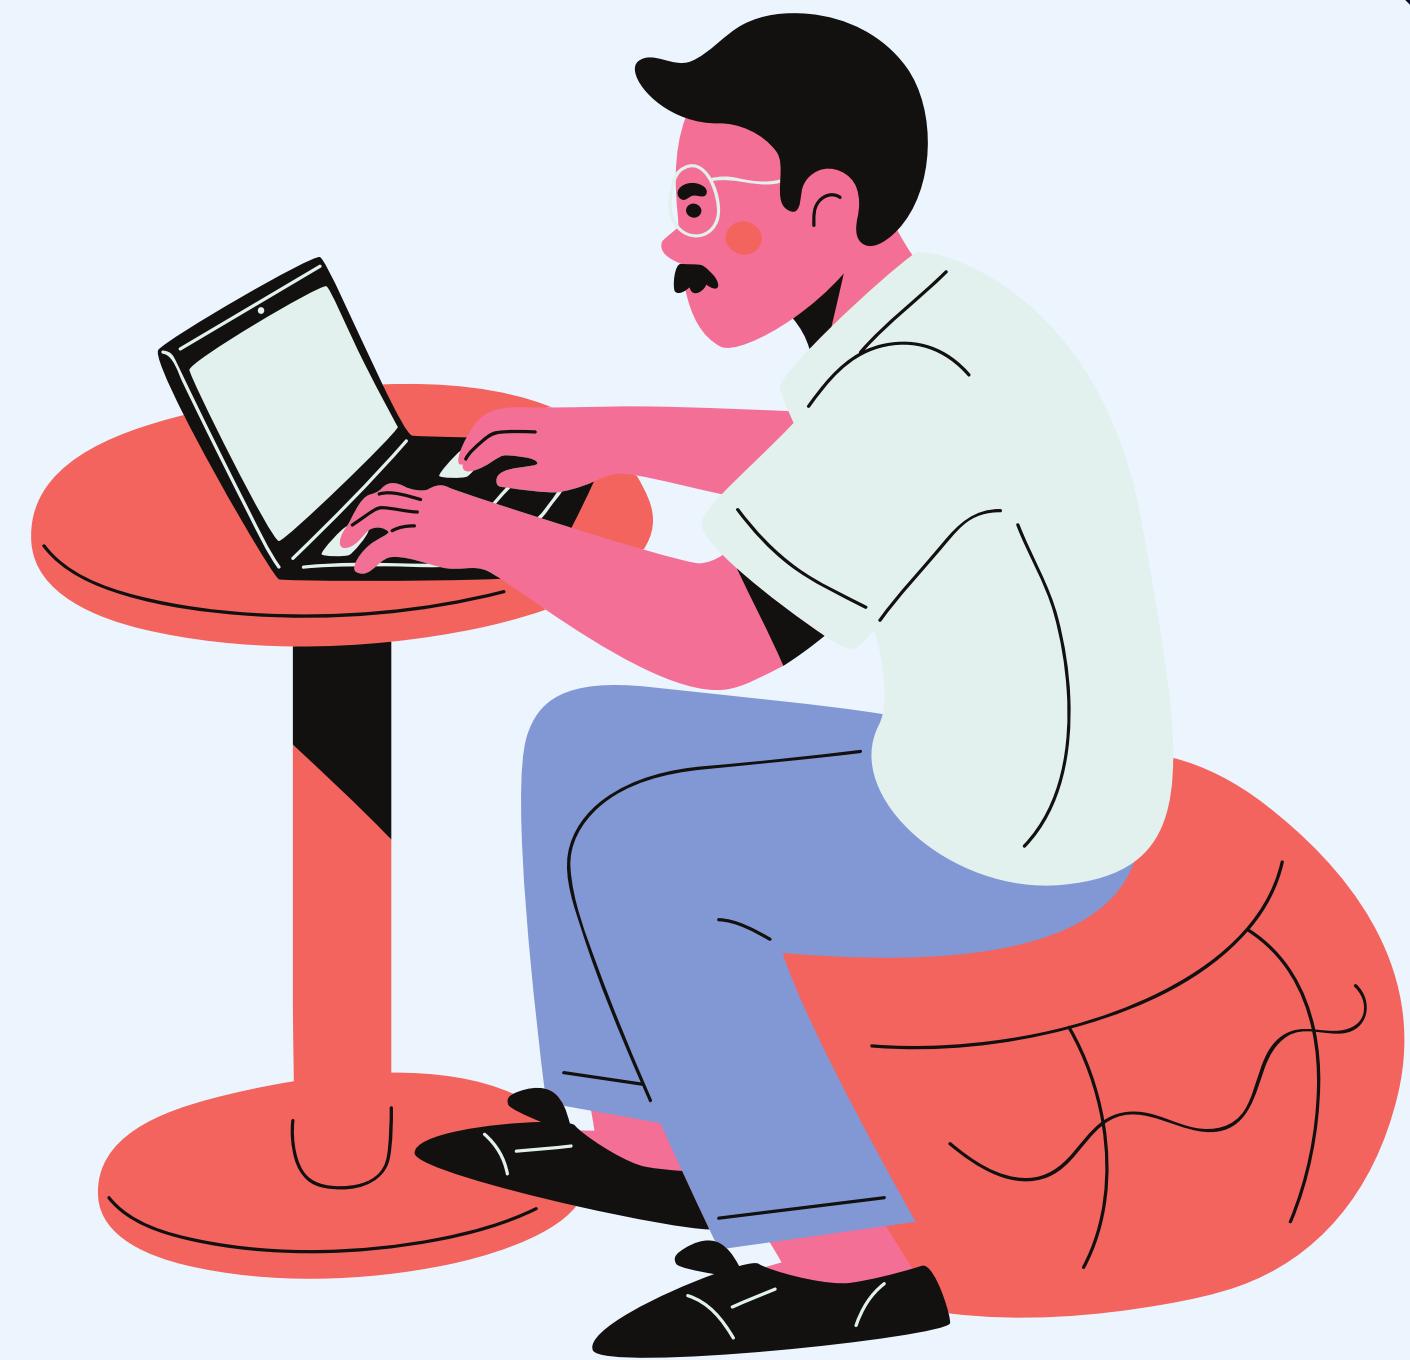


## Summary

- Mongoose query middlewares: Hooks into query lifecycle
- Pre/post-processing for validation, logging, caching
- Powerful for augmenting database operation behavior



# DEMO TIME





# AGGREGATE MIDDLEWARES



mongoose

MONGOOSE MIDDLEWARES



# AGGREGATE MIDDLEWARES



## High-Level Explanation

- Mongoose's Aggregate Middlewares: Pre and post hooks
- Intervene in aggregation pipeline execution
- Modify pipeline or process aggregated data



## Analogue

- Analogy: Aggregation pipeline as a party playlist
- `pre('aggregate')`: Adjust before the party, add excitement.
- `post('aggregate')`: Review after the party, note dance hits.



## Best Practices

- Pre-aggregate: Avoid unintentional pipeline changes.
- Post-aggregate: Watch for increased computational complexity.
- Use both judiciously to maintain code readability.



## Deep Dive

- Mongoose schema allows defining middlewares for various operations.
- `pre('aggregate')`: Executed before aggregation, modifies pipeline.
- `post('aggregate')`: Executed after aggregation, post-processing.



## When to use?

- Pre-aggregate: Enforce security policies or conditions.
  - e.g., User can only aggregate their data.
- Post-aggregate: Log for auditing, transform aggregated data.
  - e.g., Activity logging or additional data processing.

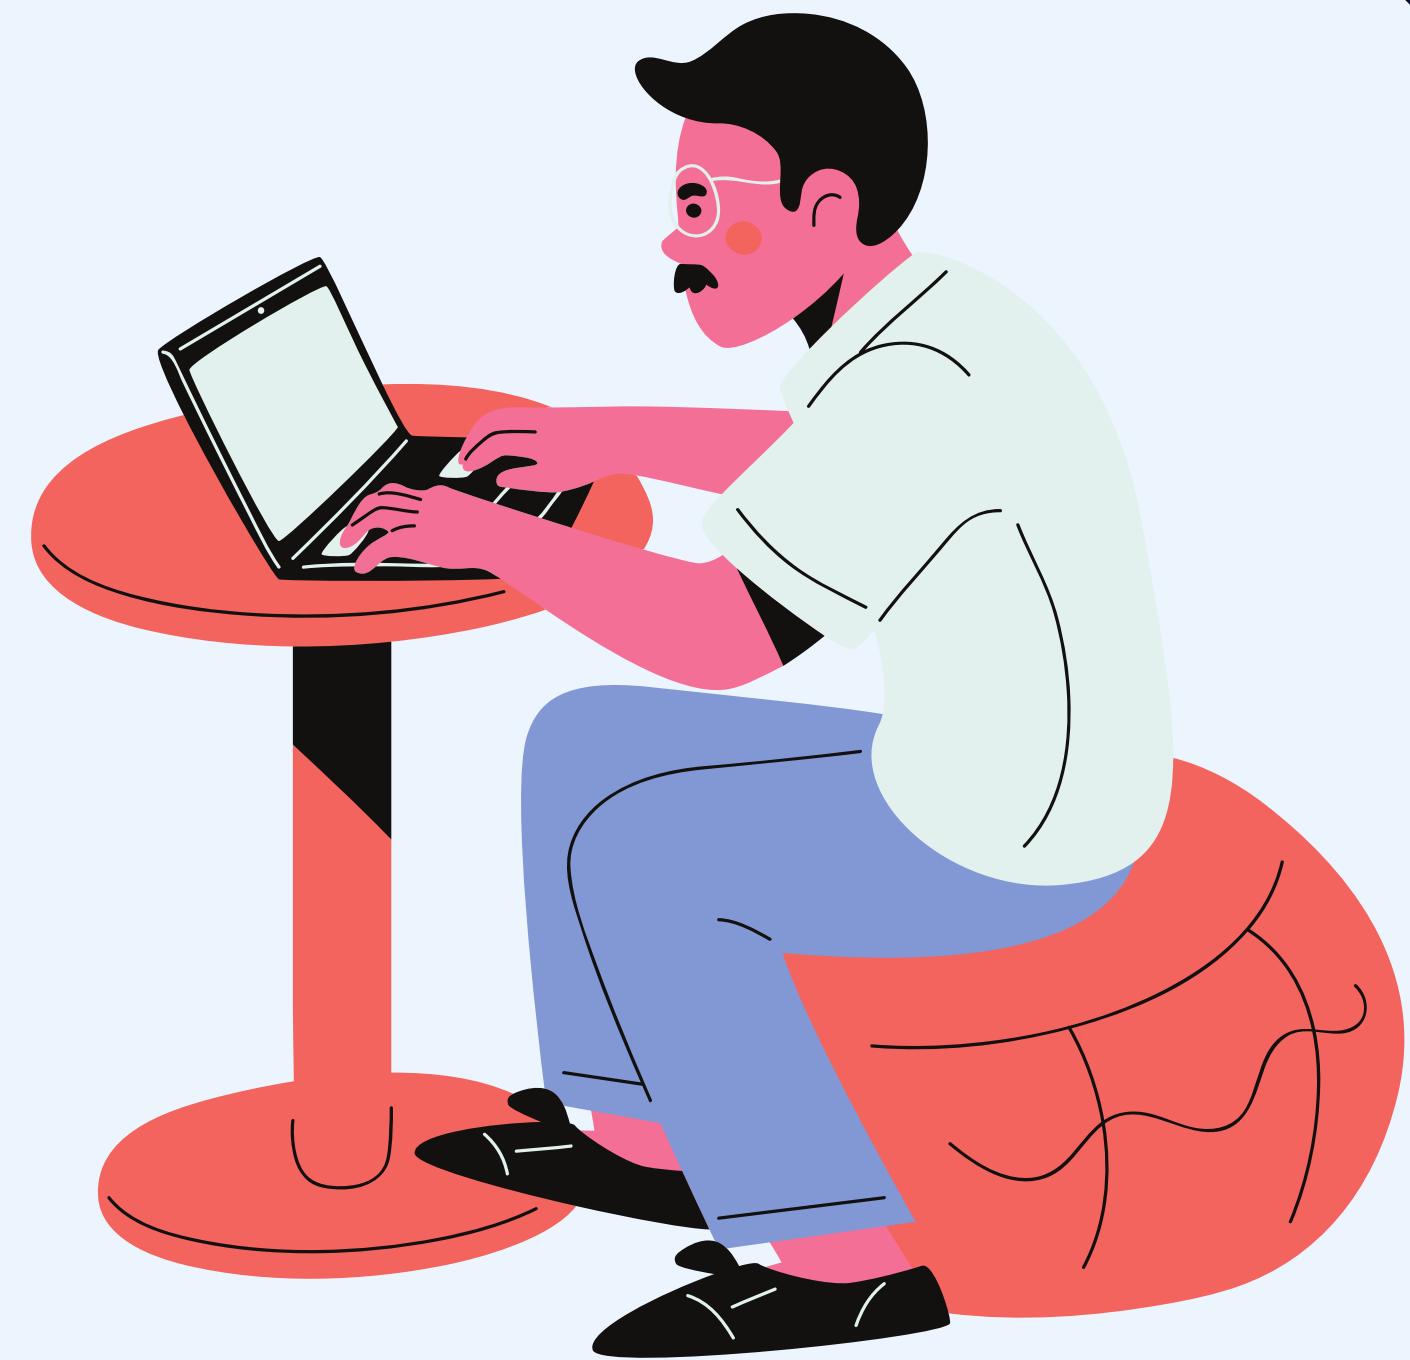


## Summary

- Mongoose Aggregate Middlewares: Hooks for pre/post aggregation.
- Useful for dynamic pipeline modification, rule enforcement, data processing.
- Use with care for code readability and performance.



# DEMO TIME





# MODEL MIDDLEWARES



mongoose

MONGOOSE MIDDLEWARES



# MODEL MIDDLEWARES



## High-Level Explanation

- Mongoose model middlewares: Pre and post-hooks for model methods.
- **Pre-hooks:** Triggered before method execution, e.g., `pre('save')`.
- **Post-hooks:** Run after method execution, handle operation results, e.g., `post('save')`.



## Analogue

- Analogy: Blog post on a website and Mongoose model middlewares
- `pre-save` hook: Spell-check or tag addition before "Publish."
- `post-save` hook: Share on social media after publication.



## Best Practices

- **Ordering:** Sequence of hook definitions matters.
- **Asynchronous Operations:** Handle async operations carefully.
- **Error Handling:** Ensure error handling, especially in pre-hooks for data validation.



## Deep Dive

- Mongoose schema allows defining middlewares for various operations.
- `pre('aggregate')`: Executed before aggregation, modifies pipeline.
- `post('aggregate')`: Executed after aggregation, post-processing.



## When to use?

- Pre-hooks: Validation, data transformation, business logic.
- Post-hooks: Notifications, logging, post-operation actions.

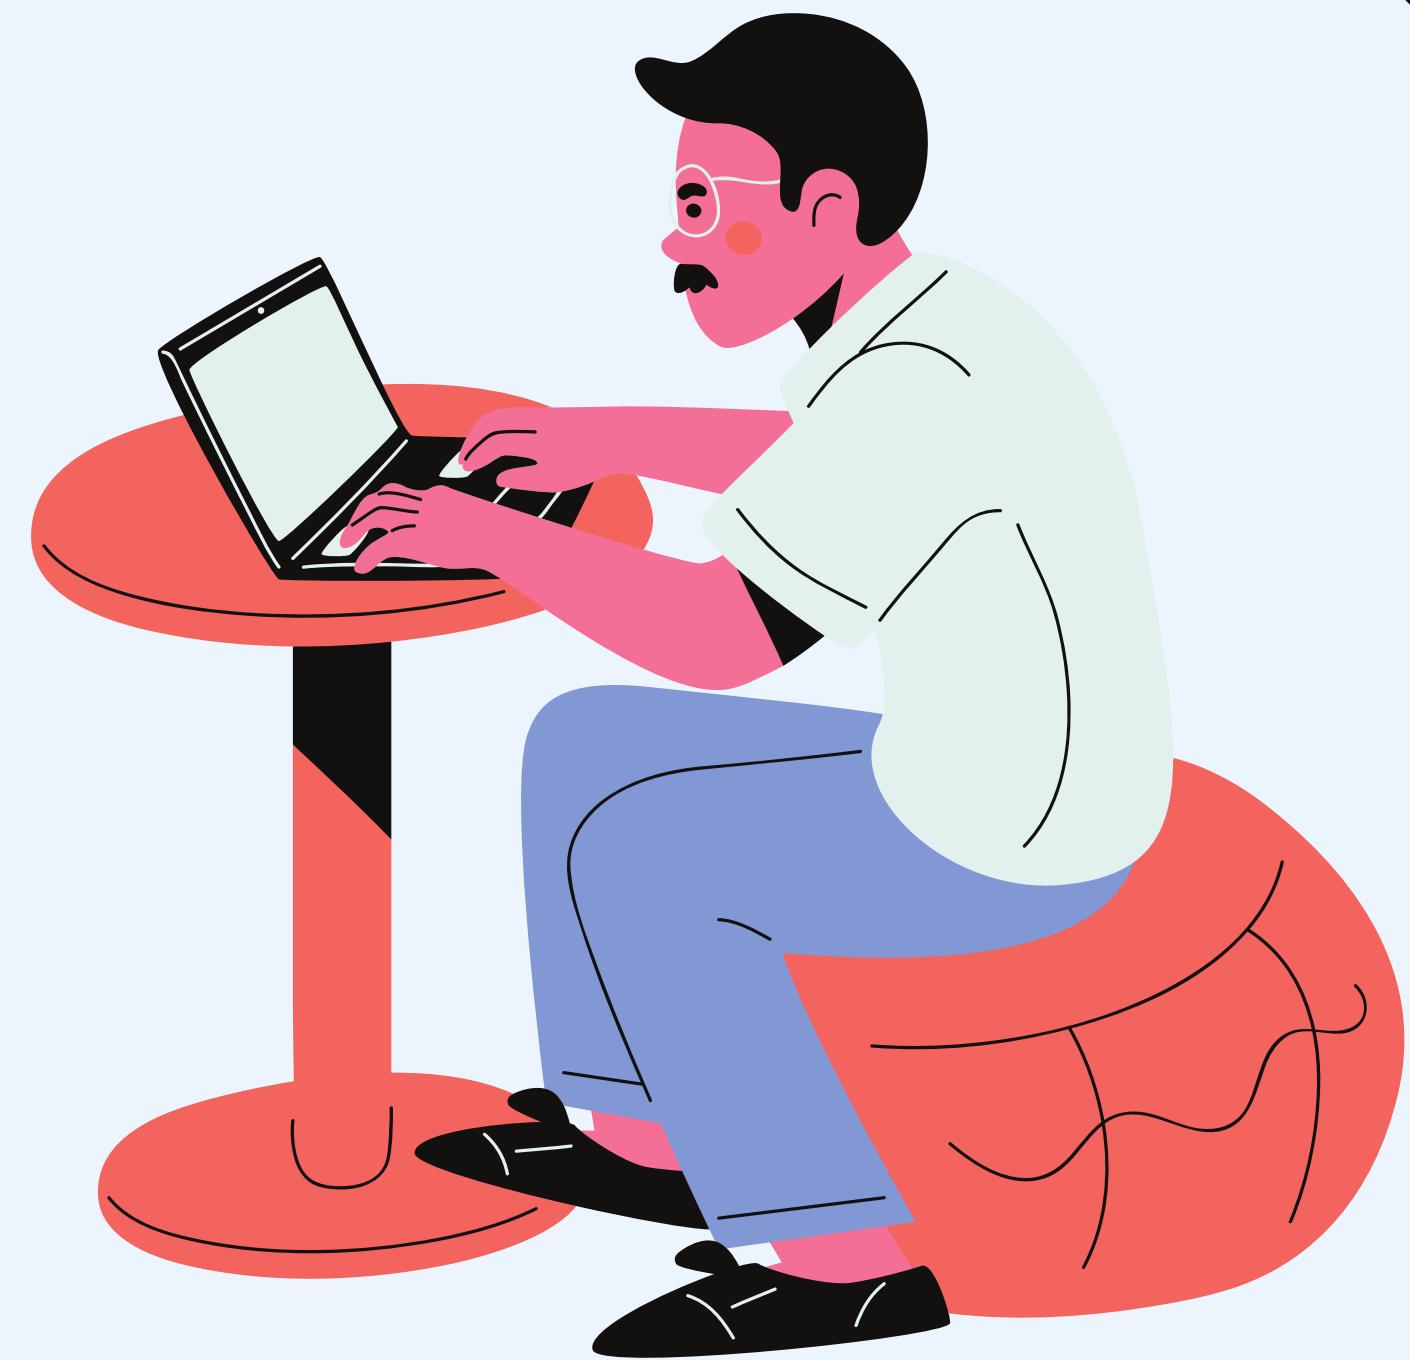


## Summary

- Mongoose Model Middlewares: Gatekeepers and observers.
- Useful for validation, transformation, logging, notifications.
- Implement with care for a clean, efficient codebase.



# DEMO TIME





# MIDDLEWARES STAGES



mongoose

MONGOOSE MIDDLEWARES



# MIDDLEWARE STAGES



## High-Level Explanation

- **init**: Data manipulation after fetch but before use.
- **validate**: Last chance data validation before save.
- **save**: Pre/post save logic, widely used.
- **remove**: Pre-document removal logic, e.g., cascading deletes, notifications.



## Analogue

- Analogy for Mongoose middleware stages:
- `init`: Starting a new game level, stage is set.
- `validate`: Checking inventory before proceeding.
- `save`: Reaching a checkpoint, progress is saved.
- `remove`: Exiting with an option to save friend's scores.



## Best Practices

- Asynchrony: Call `next()` after async operations complete.
- Error Handling: Pass errors to `next()` for proper handling.



## Deep Dive

- Mongoose schema allows defining middlewares for various operations.
- `pre('aggregate')`: Executed before aggregation, modifies pipeline.
- `post('aggregate')`: Executed after aggregation, post-processing.



## When to use?

- `init`: Initial state setup, document transformations.
- `validate`: Custom validation rules.
- `save`: Timestamps, encryption, business logic.
- `remove`: Cleanup tasks, associated record deletion.



## Summary

- Mongoose middleware stages offer precise document lifecycle control.
- Improve data integrity and application logic.
- Follow best practices: error handling, async operation management.



# DEMO TIME





# POPULATING REFERENCES AUTOMATICALLY



mongoose

MONGOOSE MIDDLEWARES

# POPULATING REFERENCES AUTOMATICALLY



## High-Level Explanation

- Mongoose Query Middleware: Populate references automatically.
- Fill referenced documents in schema before returning data to app.



## Deep Dive

- Mongoose Schema: Define fields as ObjectIds referencing other collections.
- Normally, use ` `.populate()` method to fill references.
- Query Middleware automates this, populating references on query.



## Analogue

- Analogy: Auto-complete for search bar.
- Normally, separate button for full details.
- With Query Middleware, related data auto-populated.



## When to use?

- Frequent queries needing populated references.
- When client apps expect fully formed objects, server-side population.



## Summary

- Auto-populating references: Automatically fill in data.
- Resolves references between collections.
- Simplifies client-side code.
- Query Middleware automates and streamlines this process.



# DEMO TIME





# THANK YOU

